

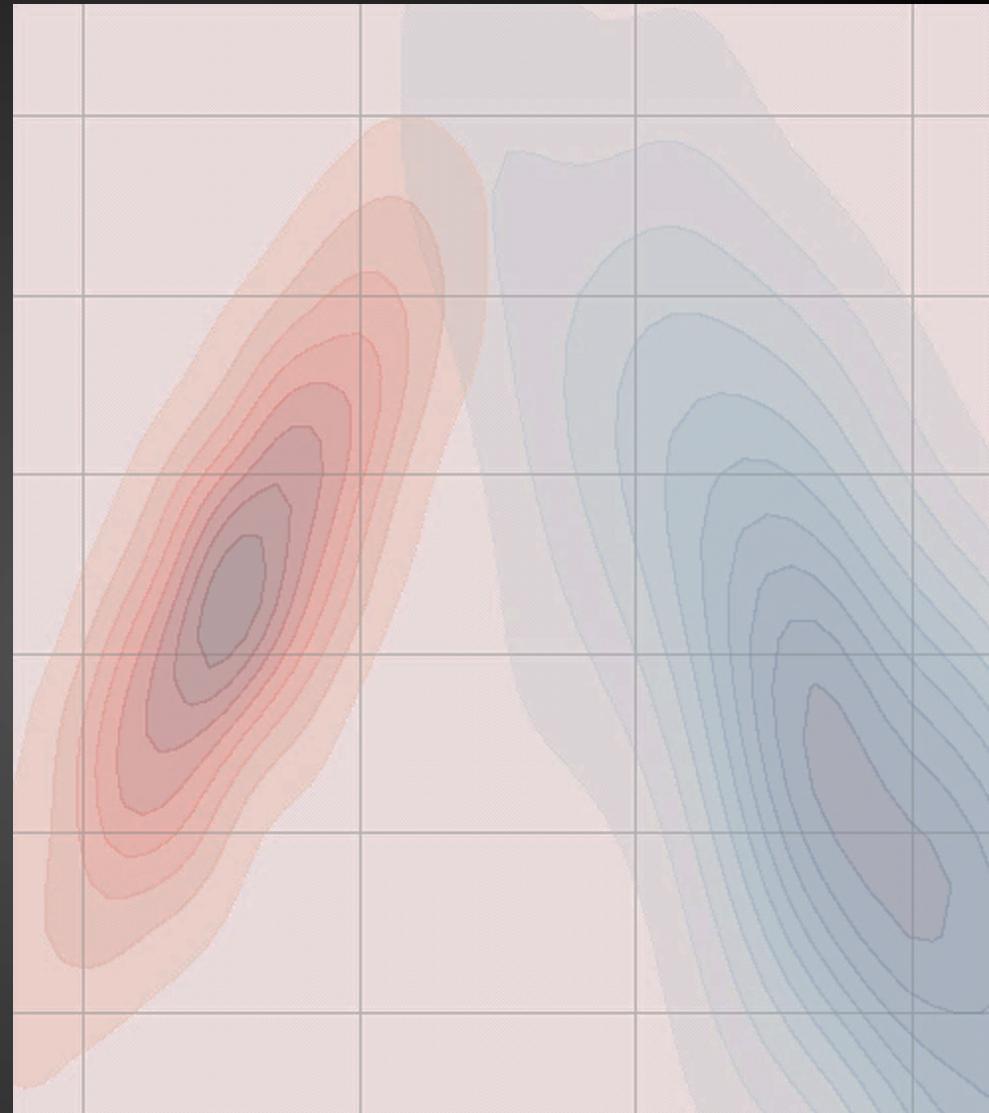
# Interrupted Time Series Experiments in Python

Drew Fustin

*Automaton Data*

[drewfustin@gmail.com](mailto:drewfustin@gmail.com)

@drewfustin



[github.com/drewfustin/talks\\_time-series-experiments](https://github.com/drewfustin/talks_time-series-experiments)

# Interrupted Time Series Experiments in Python



- Typical Experimentation in Data Science
- Time Series Experiments / Regression Discontinuity
- Simulating Time Series Data
- ARIMA Algebra and Forecasting with statsmodels
- Modeling Interruptions with PyMC3

# Me

**Automaton Data**

*available for hire!*

drew@automatondata.com

PhD, Physics

Data Scientist



SPOT  
**HERO**  
**GRUBHUB**

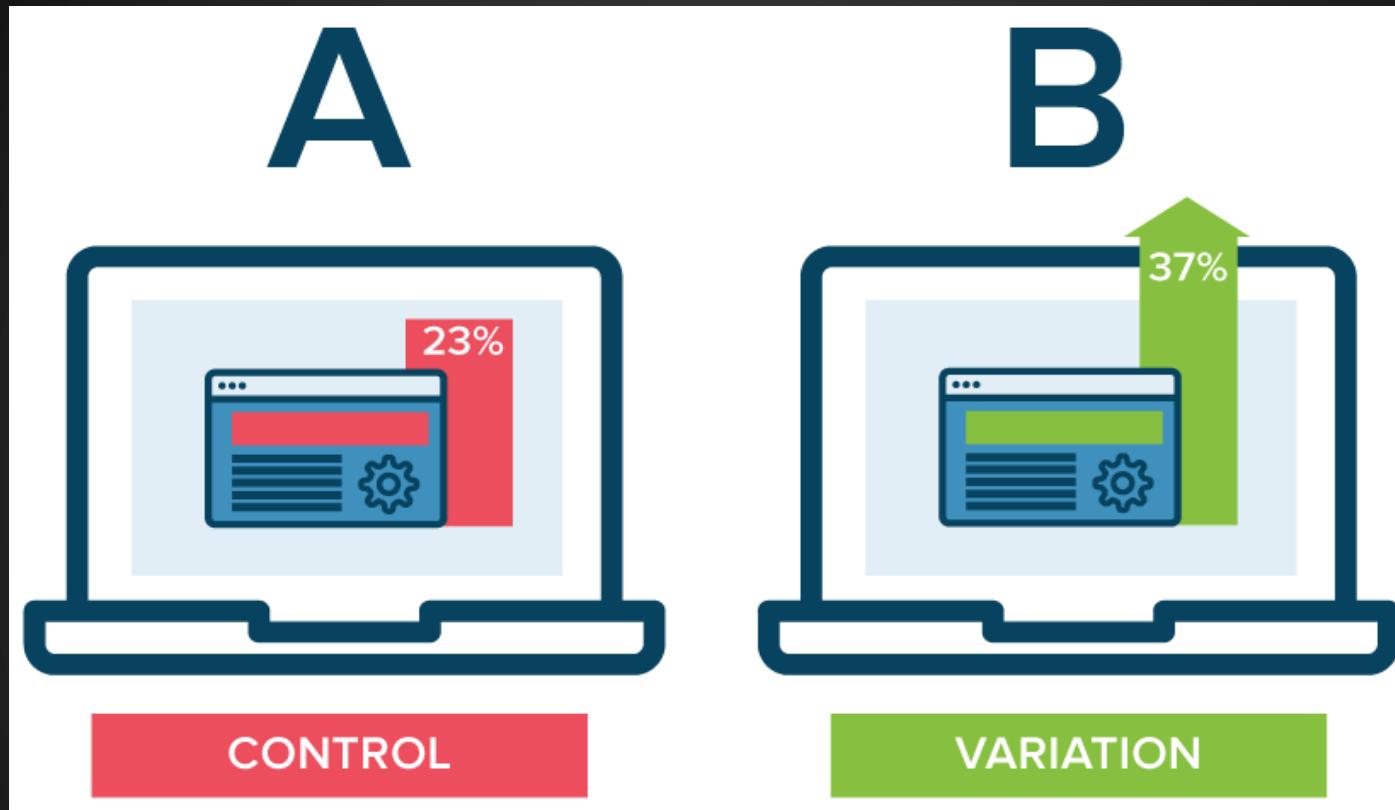


# You?



- Hypothesis Testing (e.g. A/B Testing)?
- Time Series Forecasting (e.g. ARIMA)?
- Probability Distributions (e.g. Poisson)?
- Markov Chain Monte Carlo (MCMC)?

# Typical Experiments



source: Optimizely

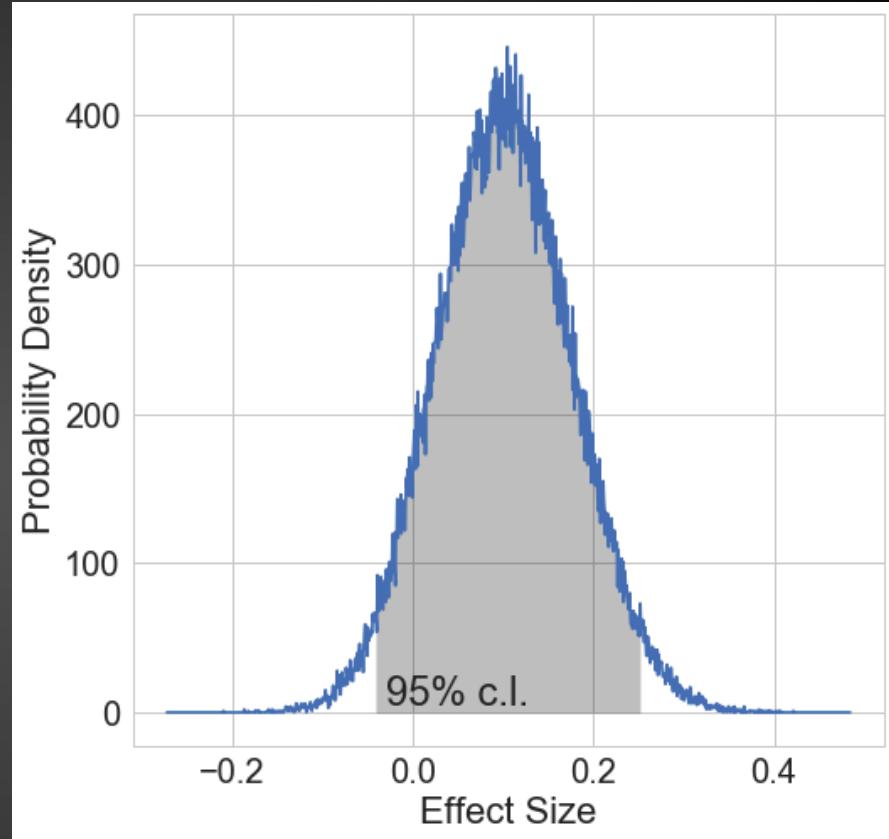
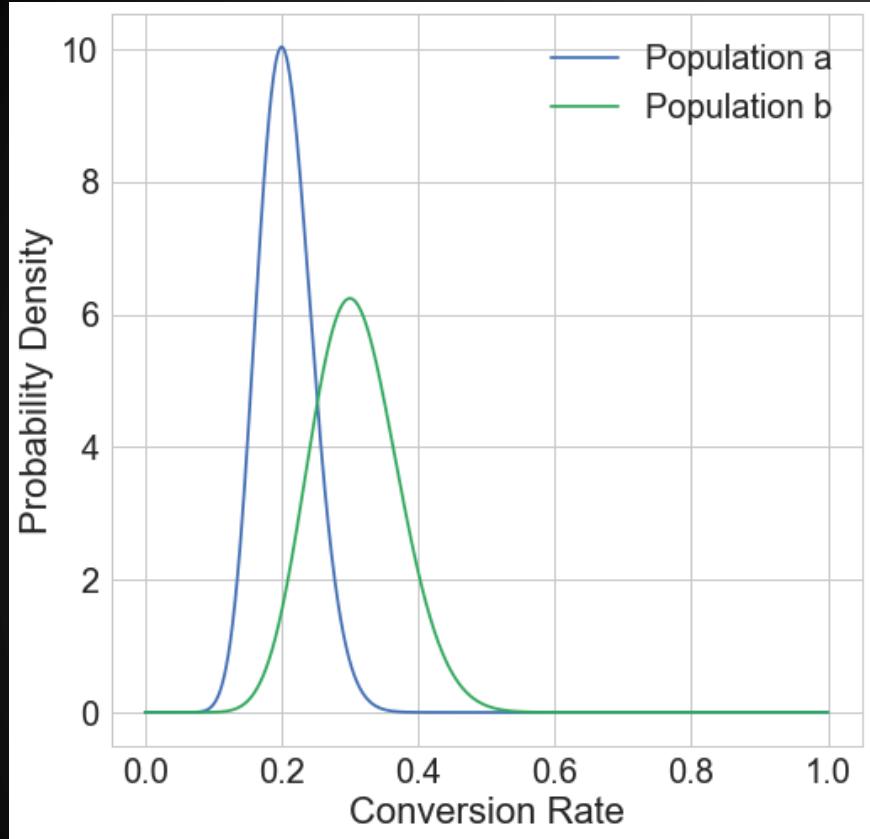
- Select Control/Variant populations (a, b)
- Vary something and measure the change
- Run concurrently to avoid unnecessary variation

# Typical Experiments

e.g. A/B testing conversion rate

```
def conversion_test(successes_a, failures_a, successes_b, failures_b):
    a_distribution = scipy.stats.beta(successes_a + 1, failures_a + 1)
    b_distribution = scipy.stats.beta(successes_b + 1, failures_b + 1)
    effect_size = b_distribution.rvs(10000) - a_distribution.rvs(10000)
    return (effect_size.mean(),
            numpy.percentile(effect_size, (1 - 0.95) / 2 * 100),
            numpy.percentile(effect_size, (1 - (1 - 0.95) / 2) * 100))
```

# Typical Experiments



a = 100 observations, 20 successes

b = 50 observations, 15 successes

Effect Size = [-4%, 25%] to 95% confidence

# Experimentation is easy!

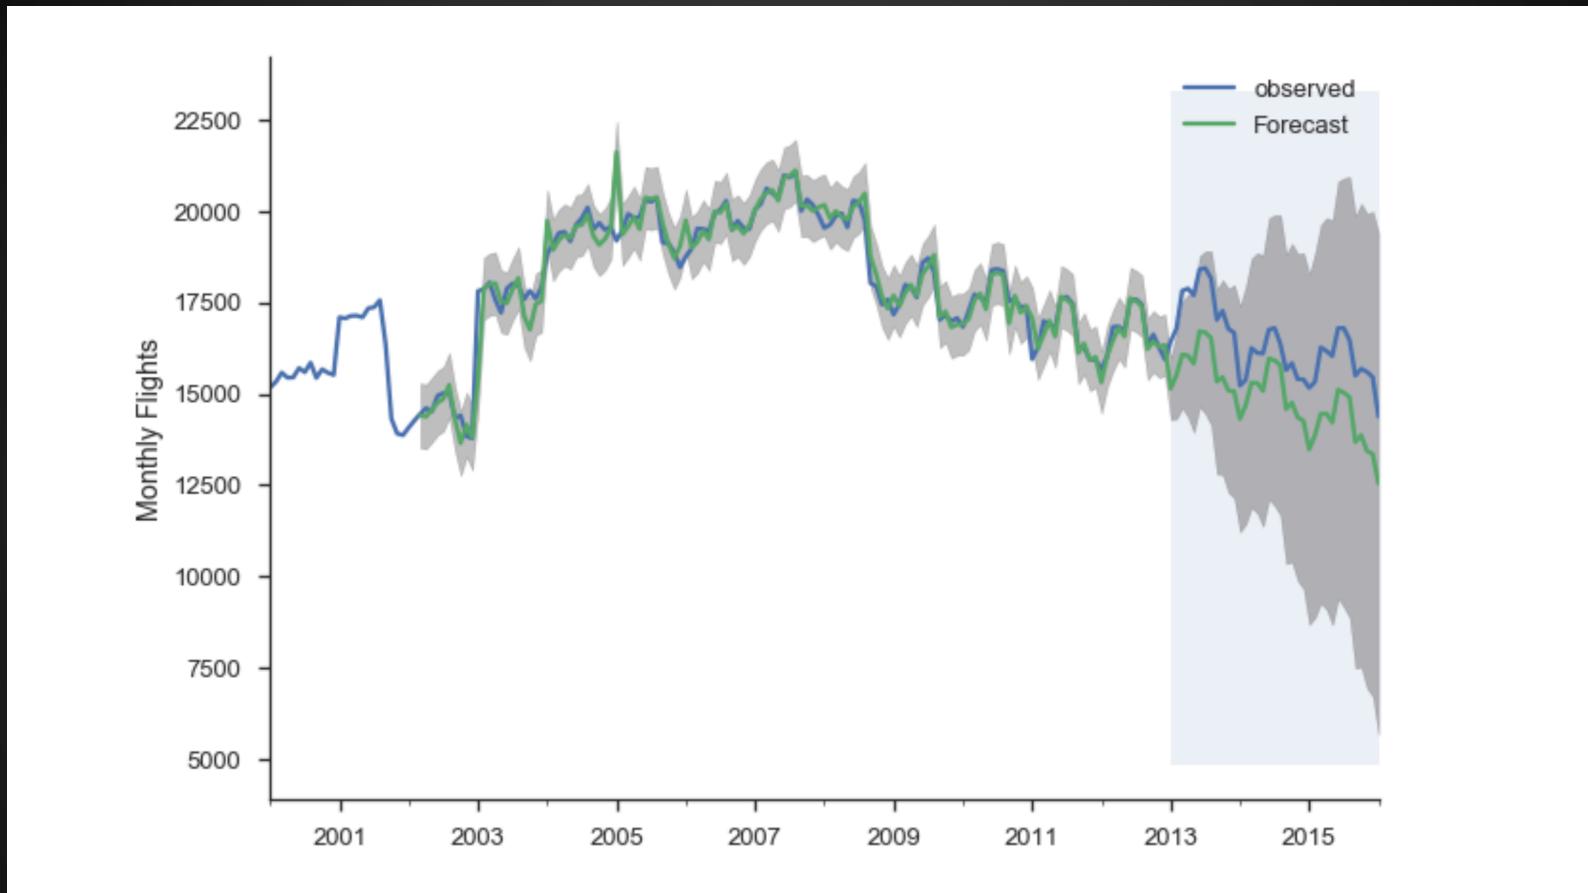


# Experimentation can be hard



I HAVE NO  
IDEA WHAT  
I'M DOING

# Experimentation can be hard



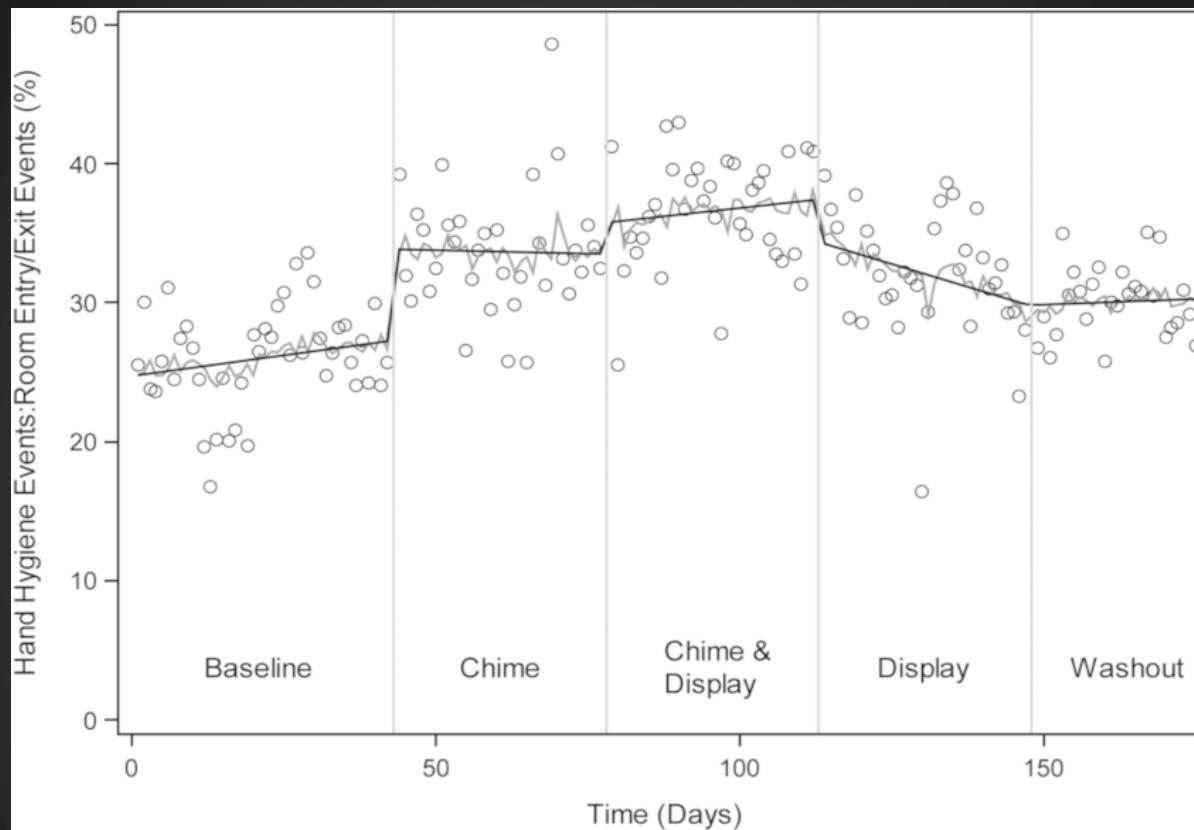
source: Tom Augspurger, [tomaugspurger.github.io/modern-7-timeseries](https://tomaugspurger.github.io/modern-7-timeseries)

What if you can't experiment on concurrent populations?

# Experimentation can be hard



# Time Series Experiments



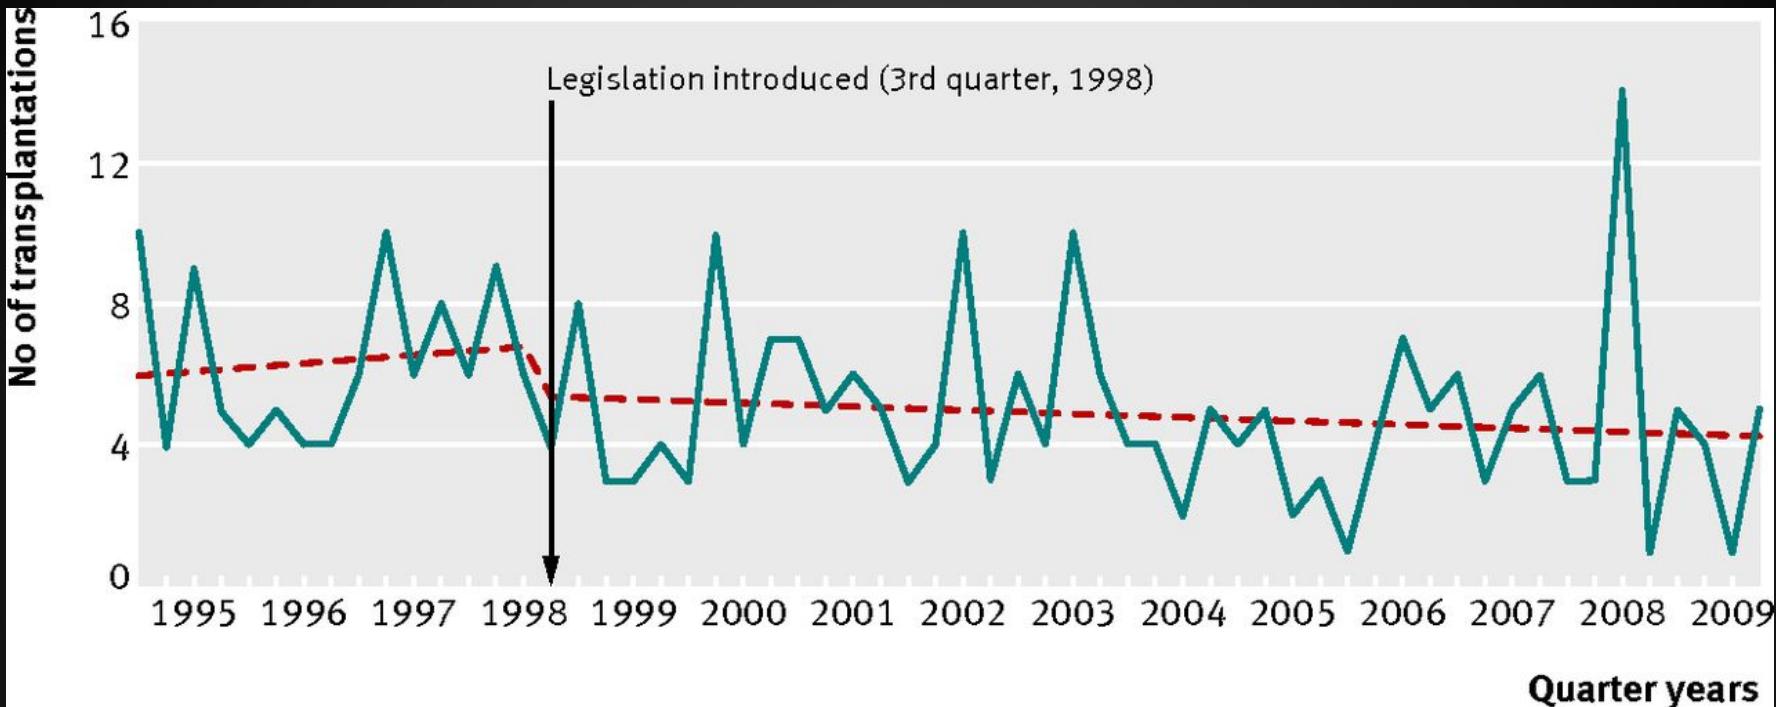
source: Ellison, et al. "A Prospective Controlled Trial of an Electronic Hand Hygiene Reminder System"

- Simply apply variation at a moment in time
- Measure before vs. after

# Time Series Experiments



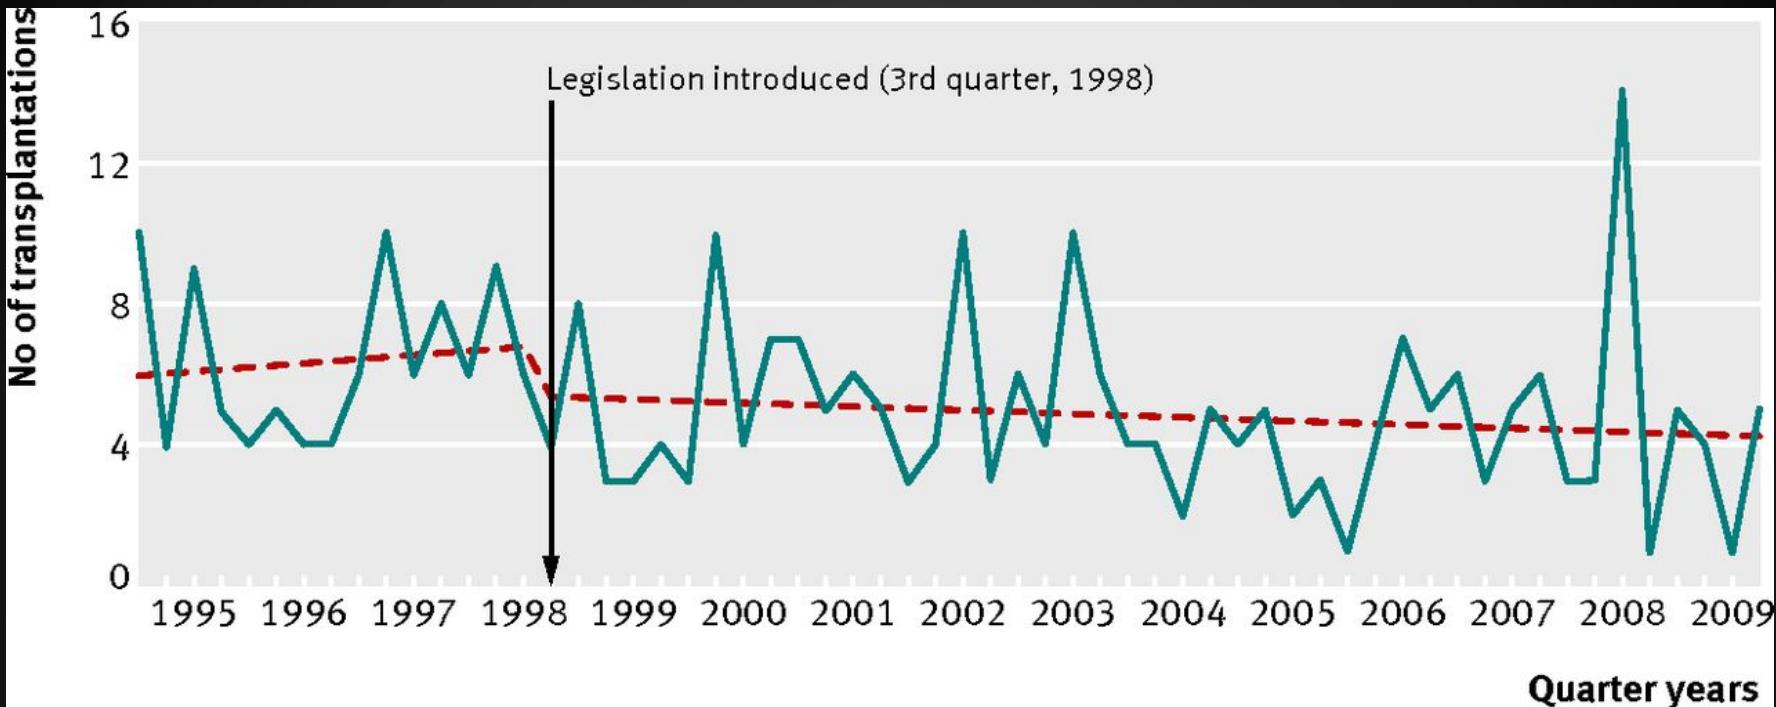
# Time Series Experiments



source: Hawton, et al. "Long term effect of reduced pack sizes of paracetamol on poisoning deaths and liver transplant activity in England and Wales: interrupted time series analyses"

- Simply apply variation at a moment in time
- Measure before vs. after

# Time Series Experiments



source: Hawton, et al. "Long term effect of reduced pack sizes of paracetamol on poisoning deaths and liver transplant activity in England and Wales: interrupted time series analyses"

- “Simply” apply variation at a moment in time
- “Measure” before vs. after

# Time Series Experiments



# Let's Give It a Try

# Make Some Data

Homogeneous Poisson Point Process

# Make Some Data

Homogeneous Poisson Point Process



points distributed randomly in time

# Make Some Data

Homogeneous Poisson Point Process



$$P(N(t) = n) = \frac{(\lambda t)^n}{n!} e^{-\lambda t}$$

$$E[N(t)] = \lambda t$$

# Make Some Data

Homogeneous Poisson Point Process

points distributed randomly in time

$$P(N(t) = n) = \frac{(\lambda t)^n}{n!} e^{-\lambda t}$$
$$E[N(t)] = \lambda t$$

$\lambda$  = constant

# Make Some Data

Homogeneous Poisson Point Process

$$P(N(t) = n) = \frac{(\lambda t)^n}{n!} e^{-\lambda t}$$
$$E[N(t)] = \lambda t$$

$\lambda$  = constant

note: the time between events that are Poisson distributed follows an **Exponential Distribution** with rate  $1/\lambda$

# Make Some Data

```
def_hpp(λ, n, t0, freq=None):
    """
    Homogeneous Poisson Point Process

    Parameters
    -----
    λ : Poisson rate parameter [transactions per day]
    n : number of transactions to produce
    t0 : initial datetime for simulation
    freq : Pandas resample rule for returned time series

    Returns
    -----
    Time series of transactions per freq
    """
    # Times between Poisson-distributed events are exponentially distributed with mean 1 / λ
    transaction_times = pd.TimedeltaIndex(np.cumsum(np.random.exponential(1 / λ, n)), unit='d')
    # Generate non-grouped HPP
    transactions = pd.Series(data=1,
                               index=pd.to_datetime(t0) + transaction_times,
                               name='transactions')
    # Group transactions by freq
    if freq:
        transactions = transactions.resample(freq).sum().fillna(0)[1:-1]

    return transactions
```

# Make Some Data

But, events never occur with a constant rate.

Can we let  $\lambda$  vary with time?

# Make Some Data

Nonhomogeneous Poisson Point Process

# Make Some Data

Nonhomogeneous Poisson Point Process

$$\lambda = \lambda(t)$$


# Make Some Data

## Nonhomogeneous Poisson Point Process

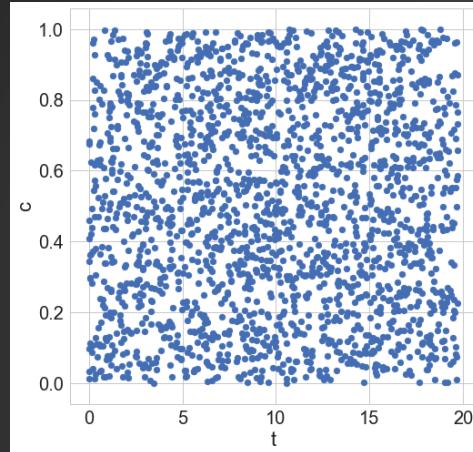
$$\lambda = \lambda(t)$$

Generation:

1. Generate HPP with rate  $\lambda_h = \max(\lambda)$
2. Attach a random number  $c_t$  bounded by  $[0, 1]$  to each point in the HPP process
3. Generate a rate function  $r(t)$  bounded by  $[0, 1]$  such that  $r(t_m) = 1$  for any  $t_m$  such that  $\lambda(t_m) = \max(\lambda)$
4. Discard any point in the HPP for which  $c_t > r(t)$

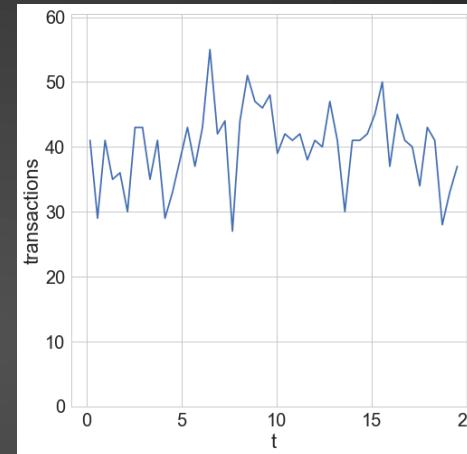
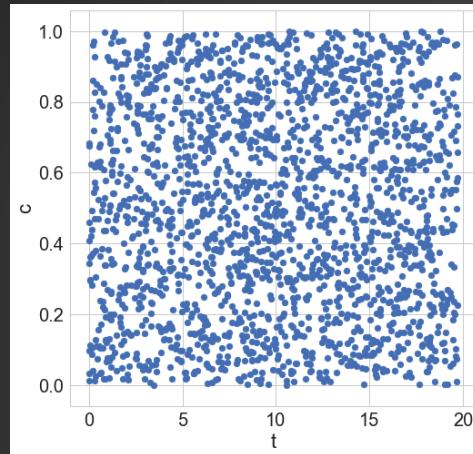
# Make Some Data

Nonhomogeneous Poisson Point Process



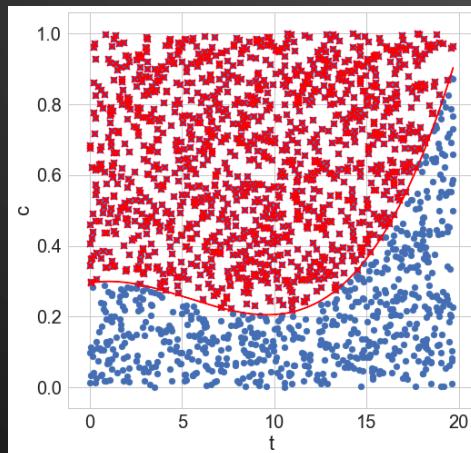
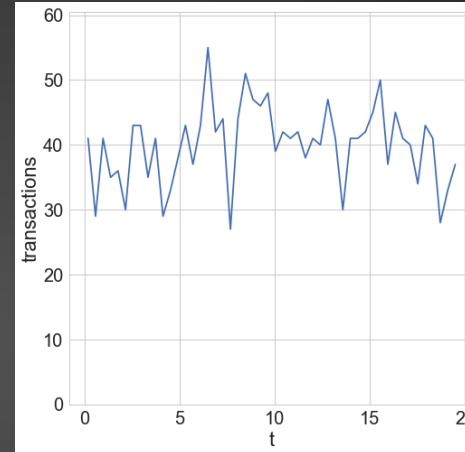
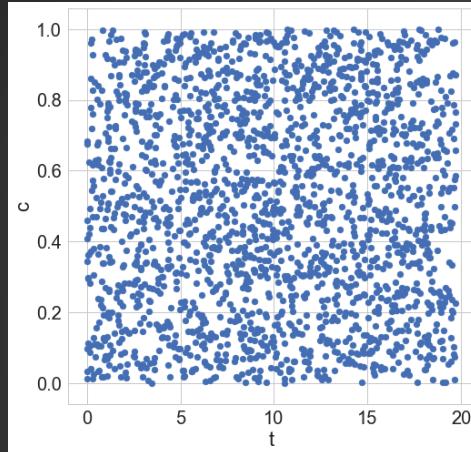
# Make Some Data

Nonhomogeneous Poisson Point Process



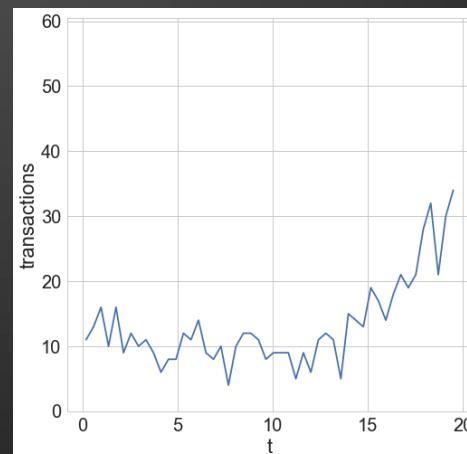
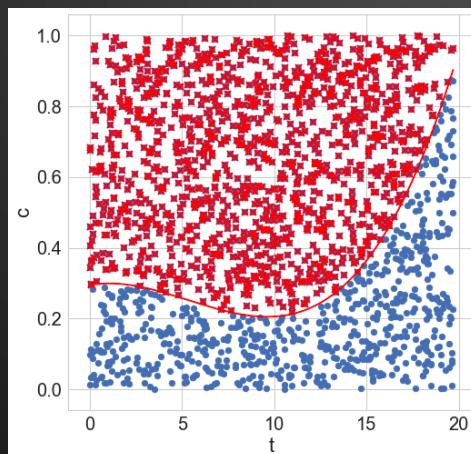
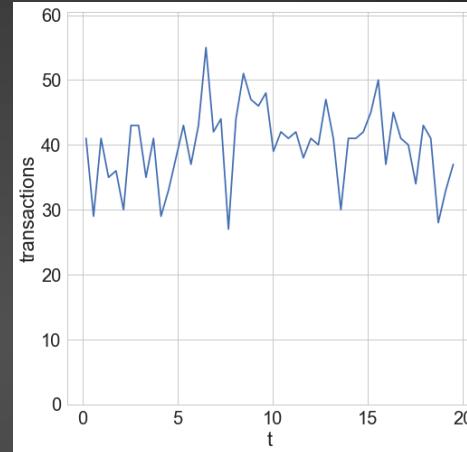
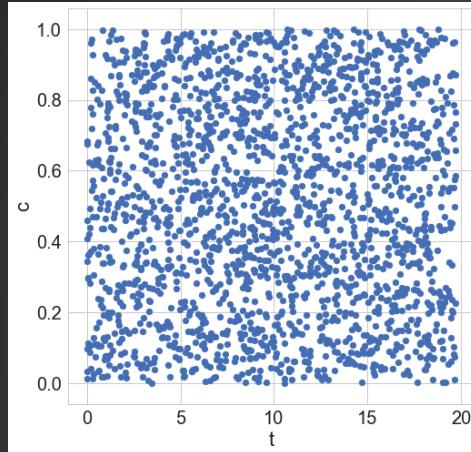
# Make Some Data

## Nonhomogeneous Poisson Point Process



# Make Some Data

## Nonhomogeneous Poisson Point Process



# Make Some Data

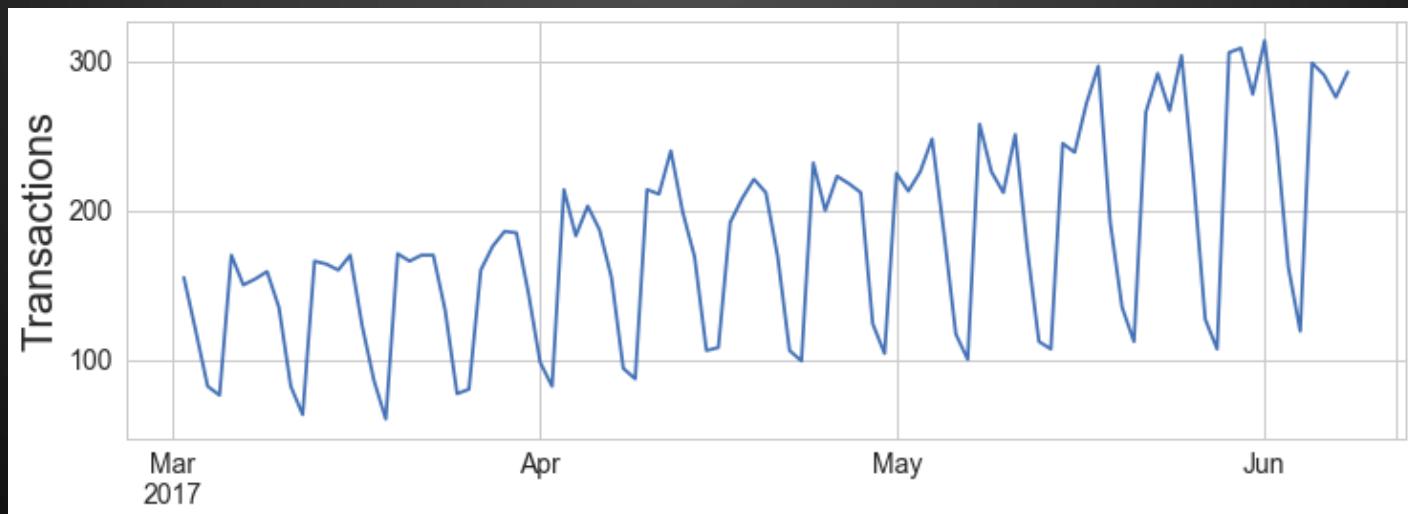
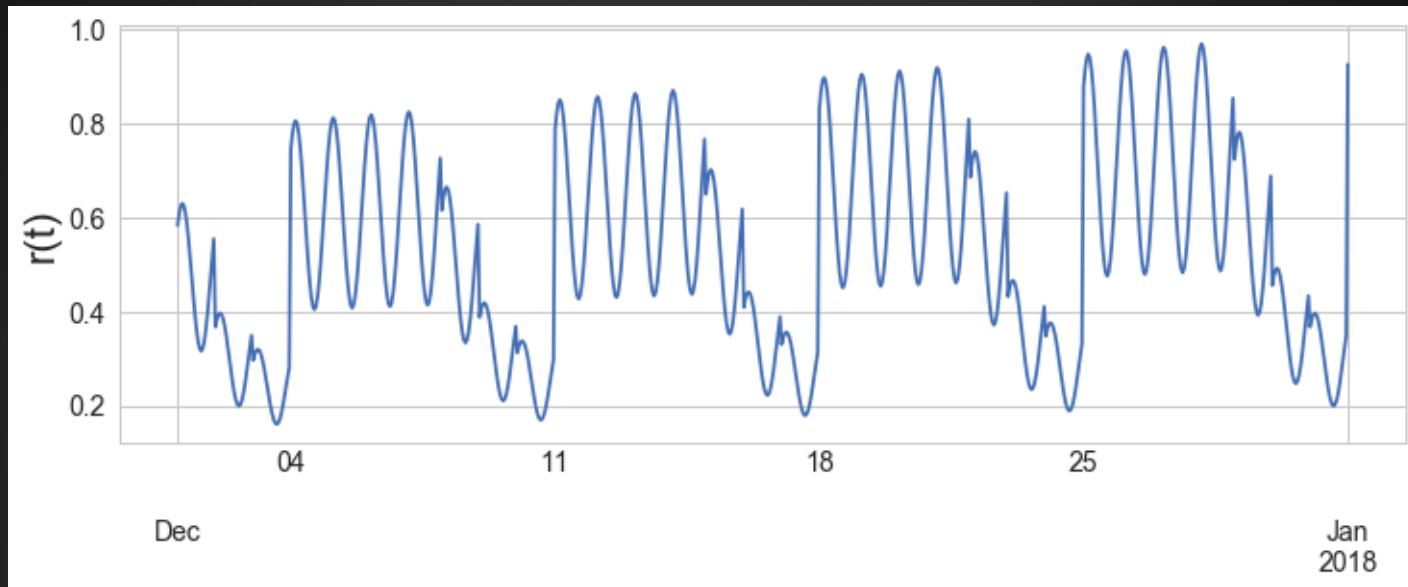
```
def nhpp(λ, n, t₀, freq=None, thinning_components=None):
    """
    Nonhomogeneous Poisson Point Process

    Parameters
    -----
    λ : Poisson rate parameter [transactions per day] of the bounding Homogeneous Point Process
        Equal to the instantaneous rate of the Nonhomogeneous Point Process when thinning = 1
    n : number of transactions to produce
    t₀ : initial datetime for simulation
    freq : Pandas resample rule for returned time series
    thinning_components : components to pass to thinning function

    Returns
    -----
    Time series of transactions per freq
    """
    # Get a non-grouped HPP
    transactions = hpp(λ, n, t₀)
    # Thin the HPP
    if thinning_components:
        transactions = transactions[np.random.uniform(size=len(transactions)) <
                                    thinning(t=transactions.index, components=thinning_components)]
    # Group transactions by freq
    if freq:
        transactions = transactions.resample(freq).sum().fillna(0)[1:-1]

    return transactions
```

# Make Some Data



# Make Some Data



# ARIMA Algebra

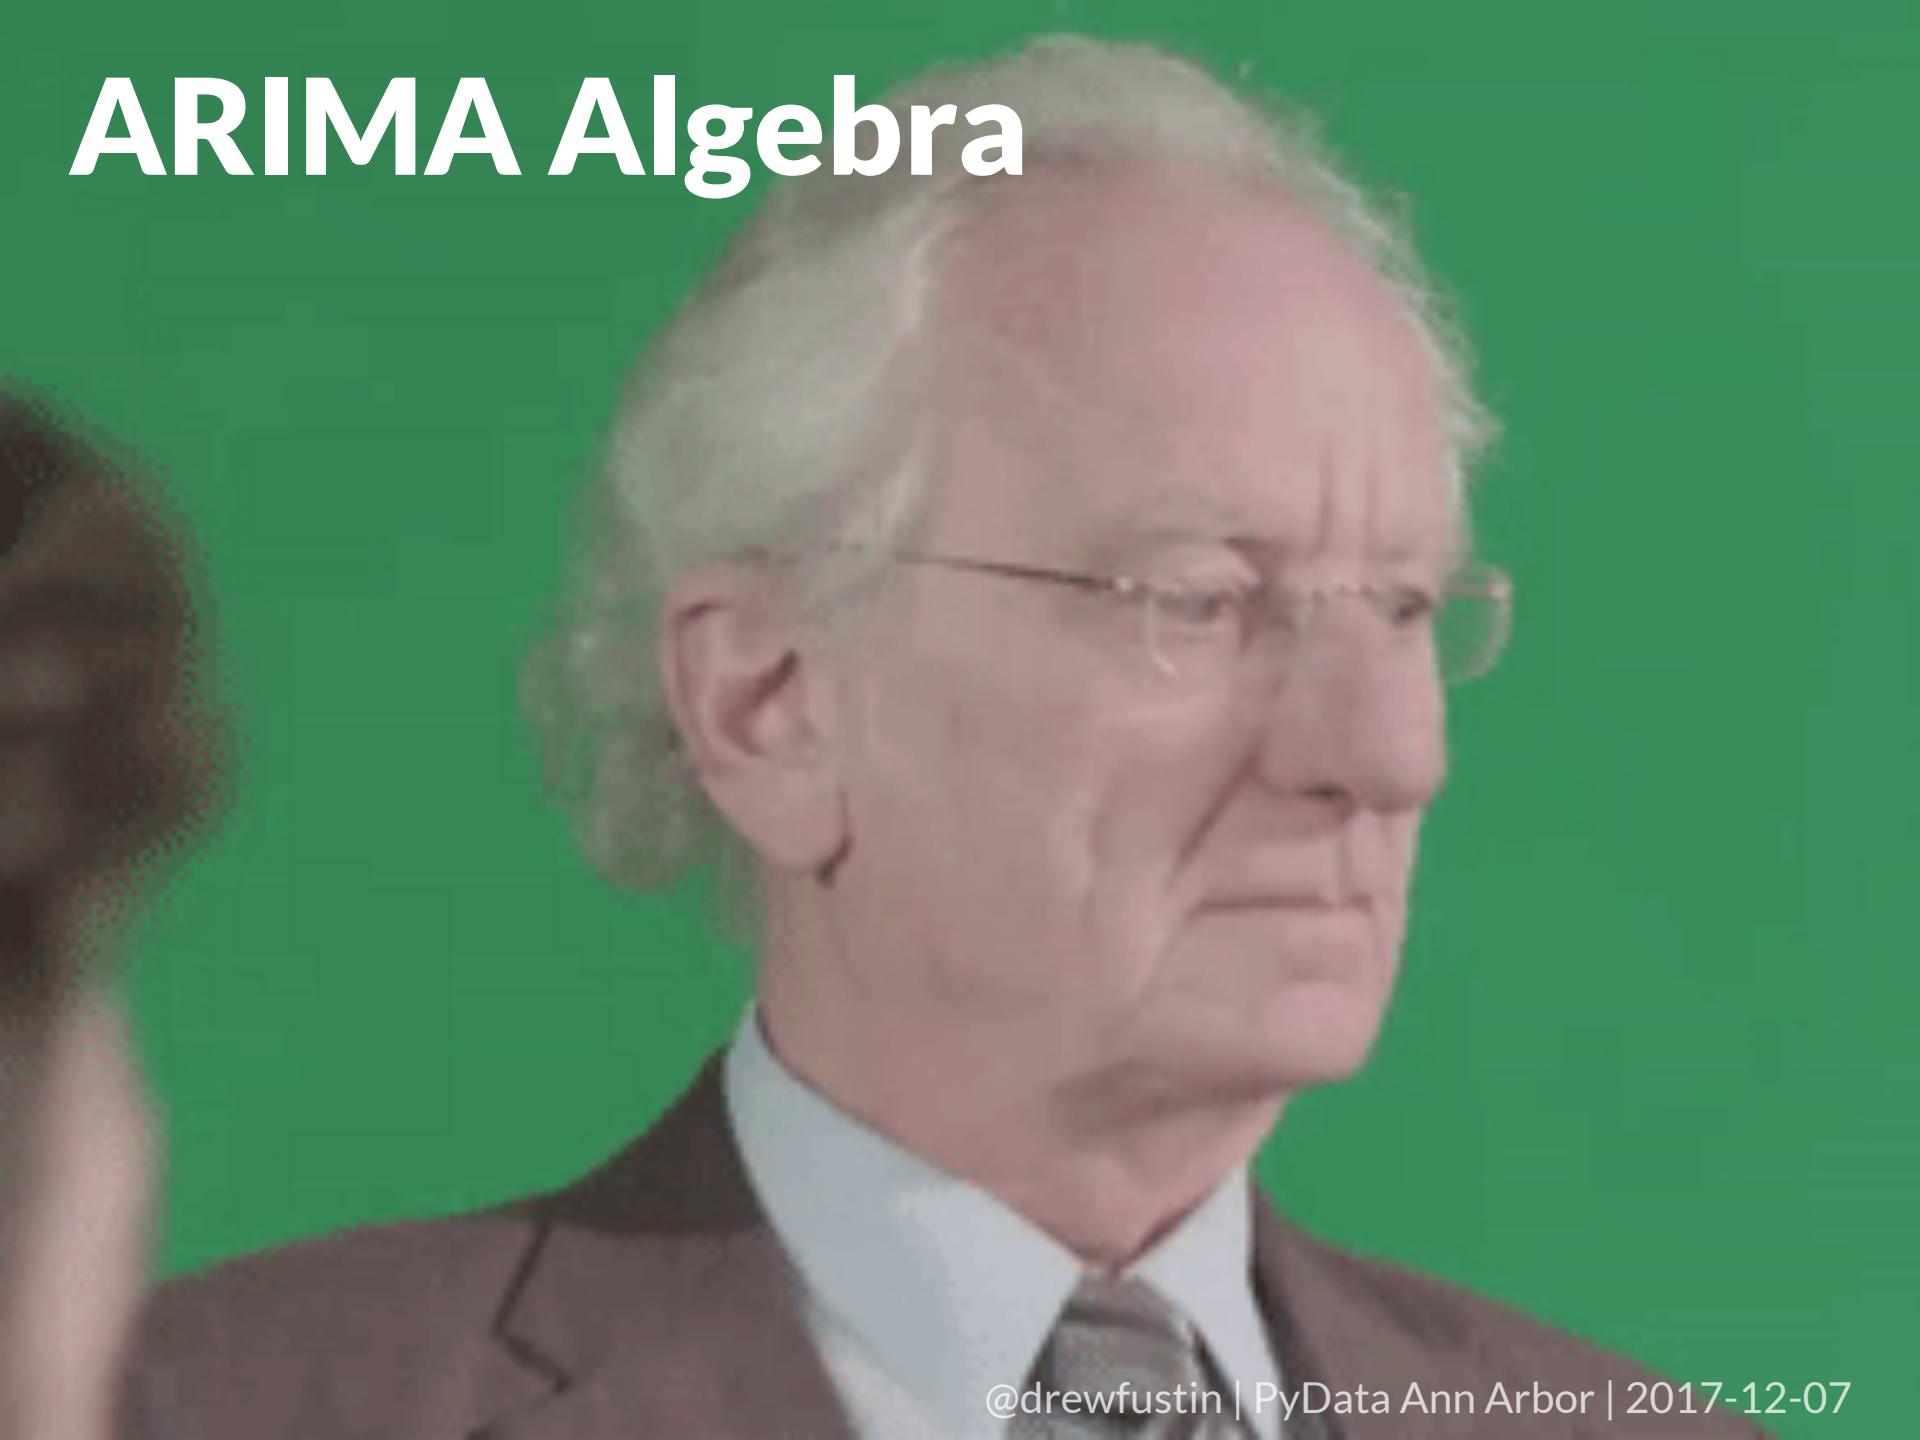
Autoregressive (**AR**) Integrated (**I**) Moving Average (**MA**)

The above can be generalized as follows.

$$\left(1 - \sum_{i=1}^p \phi_i L^i\right) (1 - L)^d X_t = \delta + \left(1 + \sum_{i=1}^q \theta_i L^i\right) \varepsilon_t.$$

This defines an ARIMA( $p,d,q$ ) process with **drift**  $\delta/(1 - \sum \phi_i)$ .

# ARIMA Algebra



# ARIMA Algebra

ARIMA ( $p$ ,  $d$ ,  $q$ )

$$\phi_p(B) \nabla_d(B) X_t = \theta_q(B) \varepsilon_t$$

# ARIMA Algebra

ARIMA ( $p, d, q$ )

$$\phi_p(B) \nabla_d(B) X_t = \theta_q(B) \varepsilon_t$$

Autoregressive (AR) [order p]

$$\phi_p(B) \equiv (1 - \phi_1 B - \phi_2 B^2 - \dots \phi_p B^p)$$

$$\phi_p(B) X_t = \varepsilon_t$$

# ARIMA Algebra

ARIMA ( $p, d, q$ )

$$\phi_p(B) \nabla_d(B) X_t = \theta_q(B) \varepsilon_t$$

Autoregressive (AR) [order p]

$$\phi_p(B) \equiv (1 - \phi_1 B - \phi_2 B^2 - \dots \phi_p B^p)$$

$$\phi_p(B) X_t = \varepsilon_t$$

Integrated (I) [order d]

$$\nabla_d(B) \equiv (1 - B)^d$$

$$\nabla_d(B) X_t = \varepsilon_t$$

# ARIMA Algebra

ARIMA ( $p, d, q$ )

$$\phi_p(B) \nabla_d(B) X_t = \theta_q(B) \varepsilon_t$$

Autoregressive (AR) [order p]

$$\phi_p(B) \equiv (1 - \phi_1 B - \phi_2 B^2 - \dots \phi_p B^p)$$

$$\phi_p(B) X_t = \varepsilon_t$$

Integrated (I) [order d]

$$\nabla_d(B) \equiv (1 - B)^d$$

$$\nabla_d(B) X_t = \varepsilon_t$$

Moving Average (MA) [order q]

$$\theta_q(B) \equiv (1 - \theta_1 B - \theta_2 B^2 - \dots \theta_q B^q)$$

$$X_t = \theta_q(B) \varepsilon_t$$

# ARIMA Algebra

This generalizes to a Seasonal ARIMA:

$$\text{ARIMA } (\textcolor{teal}{p}, \textcolor{brown}{d}, \textcolor{violet}{q}) (\textcolor{brown}{P}, \textcolor{orange}{D}, \textcolor{violet}{Q})_s \\ \phi_p(B) \Phi_P(B^s) \nabla_d(B) \nabla_D(B^s) X_t = \theta_q(B) \Theta_Q(B^s) \varepsilon_t$$

# ARIMA Algebra

This generalizes to a Seasonal ARIMA:

$$\text{ARIMA } (\textcolor{teal}{p}, d, q) (P, \textcolor{brown}{D}, \textcolor{violet}{Q})_s \\ \phi_p(B)\Phi_P(B^s)\nabla_d(B)\nabla_D(B^s)X_t = \theta_q(B)\Theta_Q(B^s)\varepsilon_t$$

e.g. Seasonal Moving Average (MA) [order Q]

$$\Theta_Q(B^s) \equiv (1 - \Theta_1 B^s - \Theta_2 B^{2s} - \dots - \Theta_Q B^{Ps})$$

# ARIMA Algebra

e.g. ARIMA (1, 1, 0):

$$\phi_1(B) \nabla_1(B) X_t = \theta_0(B) \varepsilon_t$$

$$(1 - \phi_1 B)(1 - B) X_t = (1) \varepsilon_t$$

$$(1 - \phi_1 B - B + \phi_1 B^2) X_t = \varepsilon_t$$

$$X_t - (1 + \phi_1) X_{t-1} + \phi_1 X_{t-2} = \varepsilon_t$$

# ARIMA Algebra



BBC  
AMAZON

@drewfustin | PyData Ann Arbor | 2017-12-07

# Forecasting with ARIMA

## Box-Cox Transformation

power transformation to impose normality on the time series

$$y^{(\lambda)} = \begin{cases} \frac{y^\lambda - 1}{\lambda} & \text{if } \lambda \neq 0 \\ \ln(y) & \text{if } \lambda = 0 \end{cases}$$

```
from utils import poisson_process
from scipy.stats import boxcox

comps = [
    {'function': poisson_process._trend,
     'params': {'percent_increase': 1.0, 'period': 365}},
    {'function': poisson_process._cyclical,
     'params': {'weekday_factor': [1, 1, 1, 1, 1, 0.75, 0.75]}},
    {'function': poisson_process._seasonal,
     'params': {'peak_time': '09:00:00', 'amplitude': 0.75}}]

nhpp = poisson_process.nhpp(lambda=200, n=int(1e5), t0='2017-03-01', freq='1D',
                             thinning_components=comps)
nhpp_bc, lambda_bc = boxcox(nhpp)
```

# Forecasting with ARIMA

## (Partial) Autocorrelation Functions ACF/PACF plots

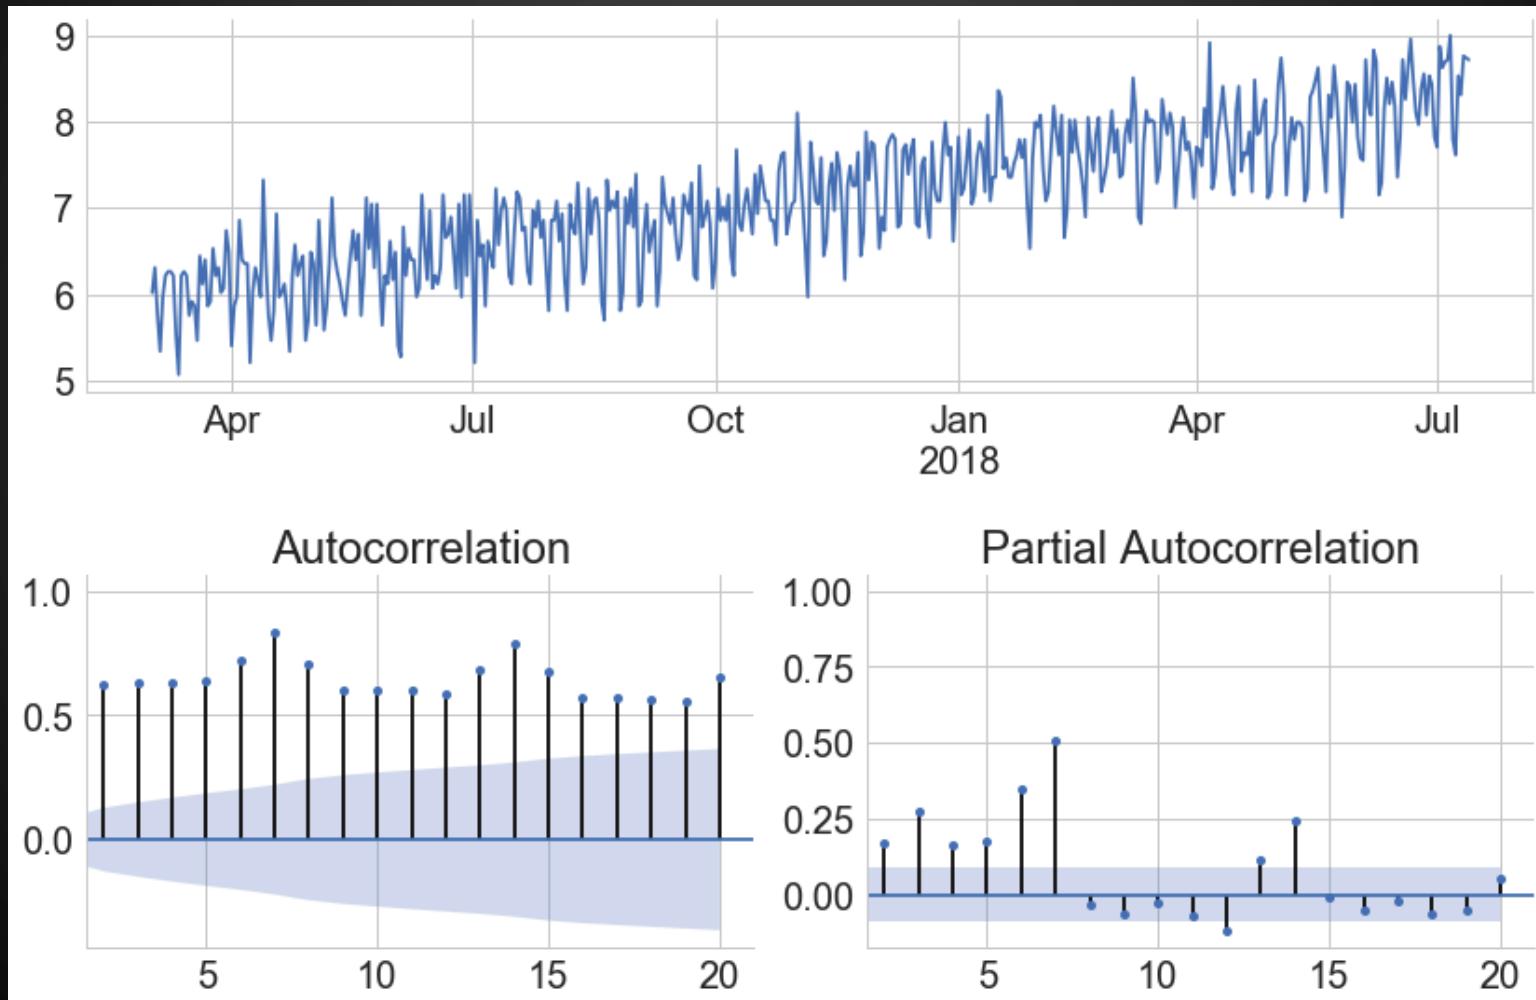
```
import matplotlib.pyplot as plt
import statsmodels.tsa.api as smt
import seaborn as sns

def tsplot(ts, lags=None):
    """
    Plot the time series with ACF and PACF
    Stolen from Tom Augspurger's fantastic post on pandas and time series forecasting
    https://tomaugspurger.github.io/modern-7-timeseries
    """
    fig = plt.figure()
    layout = (2, 2)
    ts_ax = plt.subplot2grid(layout, (0, 0), colspan=2)
    acf_ax = plt.subplot2grid(layout, (1, 0))
    pacf_ax = plt.subplot2grid(layout, (1, 1))

    ts.plot(ax=ts_ax)
    smt.graphics.plot_acf(ts, lags=lags, ax=acf_ax)
    smt.graphics.plot_pacf(ts, lags=lags, ax=pacf_ax)
    [ax.set_xlim(1.5) for ax in [acf_ax, pacf_ax]]
    sns.despine()
    plt.tight_layout()
    return ts_ax, acf_ax, pacf_ax
```

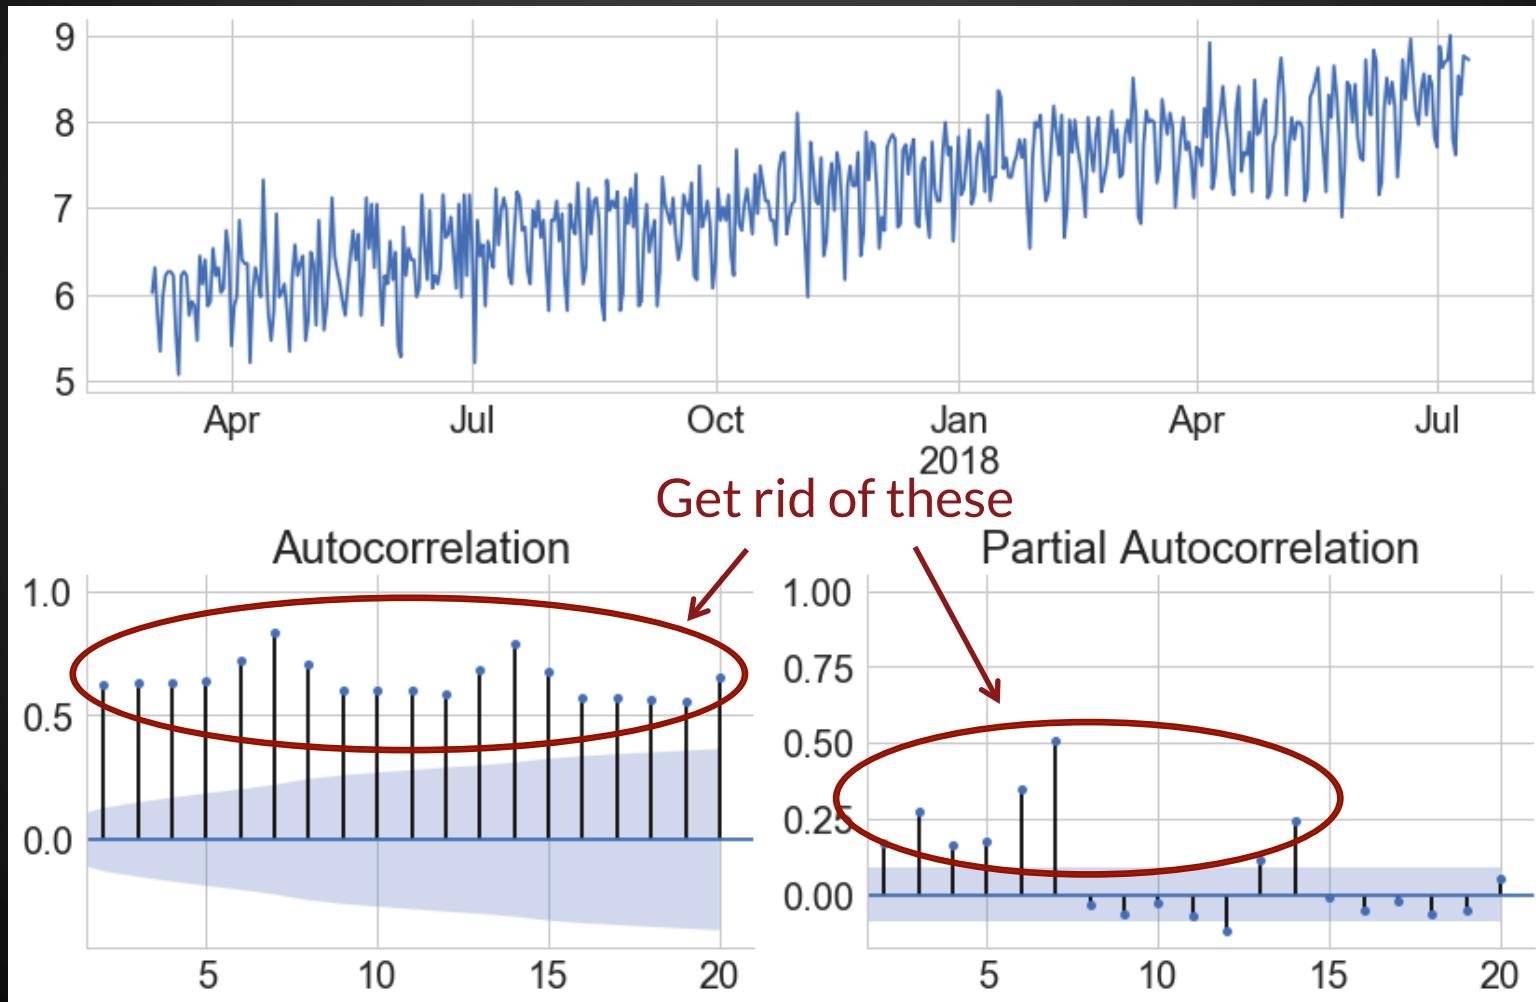
# Forecasting with ARIMA

(Partial) Autocorrelation Functions ACF/PACF plots



# Forecasting with ARIMA

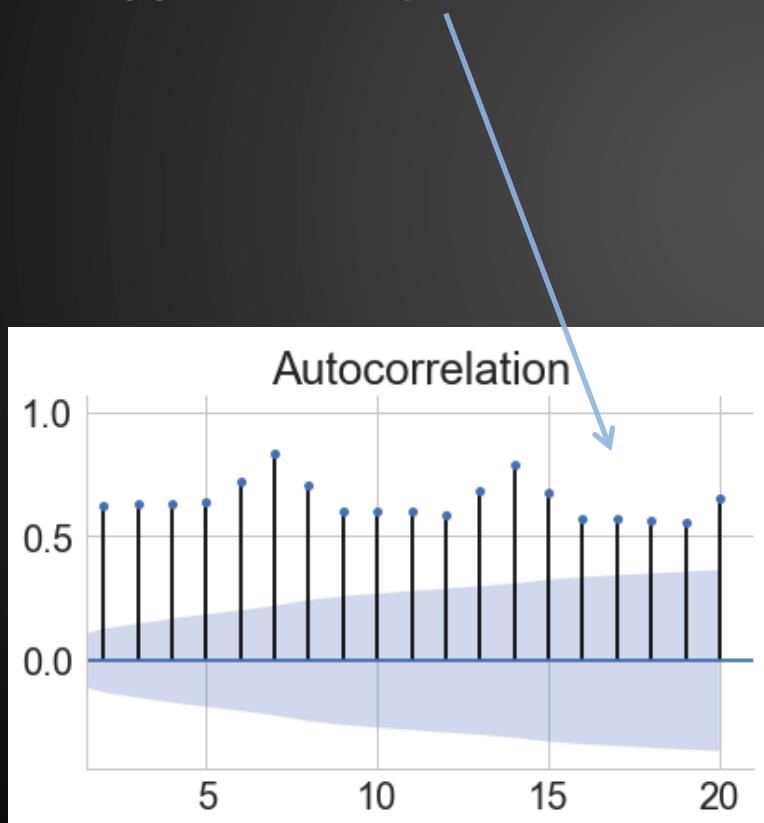
(Partial) Autocorrelation Functions ACF/PACF plots



# Forecasting with ARIMA

(Partial) Autocorrelation Functions ACF/PACF plots

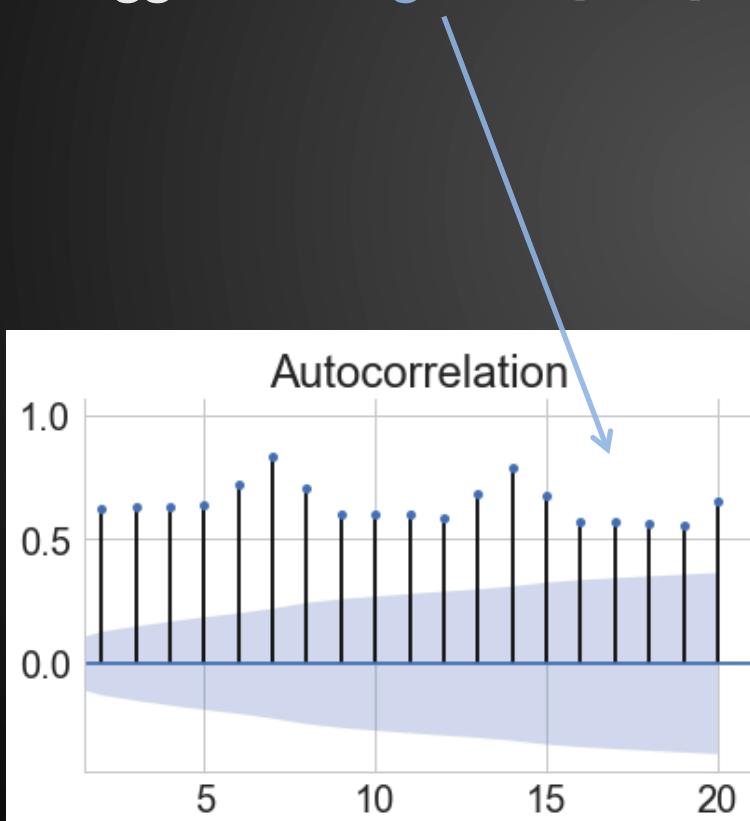
ACF being so high for so long  
suggests Integrated [ $d>0$ ]



# Forecasting with ARIMA

## (Partial) Autocorrelation Functions ACF/PACF plots

ACF being so high for so long suggests Integrated [d>0]



as does a test for Stationarity, Adjusted Dickey-Fuller Test

(this is the p-value for the time series being stationary)

```
1 import statsmodels.tsa.api as smt  
2 smt.adfuller(nhpp_bc)  
  
(-1.0652843554507179,  
 0.72871575782842646,  
 14,  
 485,  
 {'1%': -3.4439051505128342,  
 '10%': -2.5699539005207779,  
 '5%': -2.8675177321998131},  
 1083.6330094891273)
```

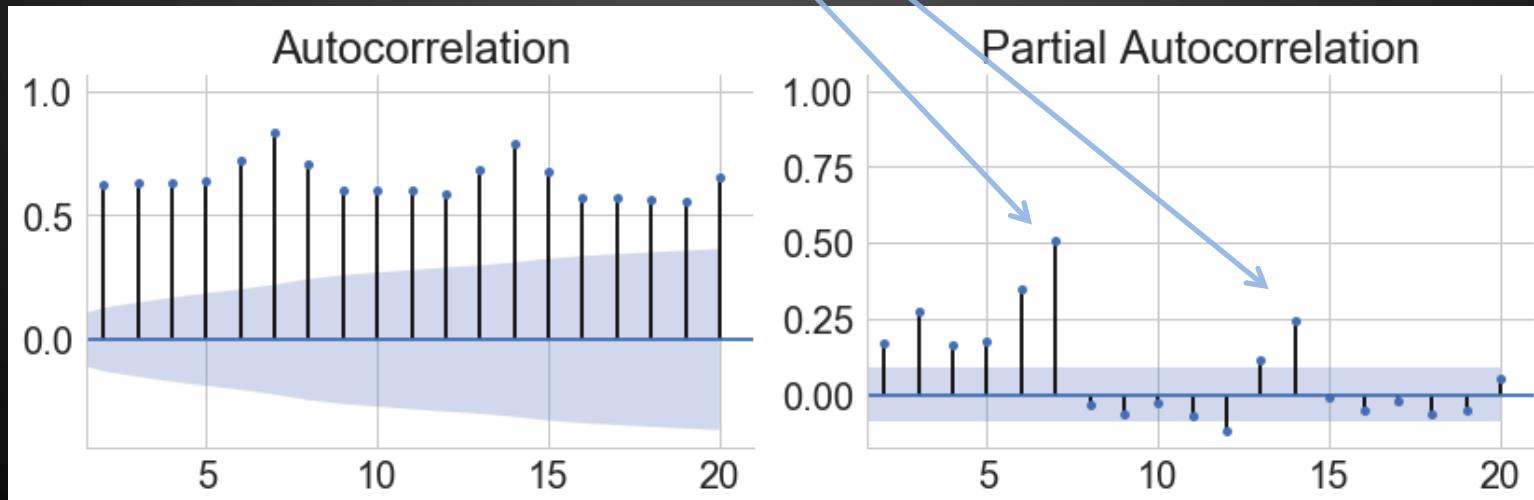
# Forecasting with ARIMA

(Partial) Autocorrelation Functions ACF/PACF plots

ACF being so high for so long  
suggests Integrated [ $d>0$ ]

suggests Seasonal [ $s=7$ ]

as does a test for Stationarity,  
Adjusted Dickey-Fuller Test



# Forecasting with ARIMA

(Partial) Autocorrelation Functions ACF/PACF plots

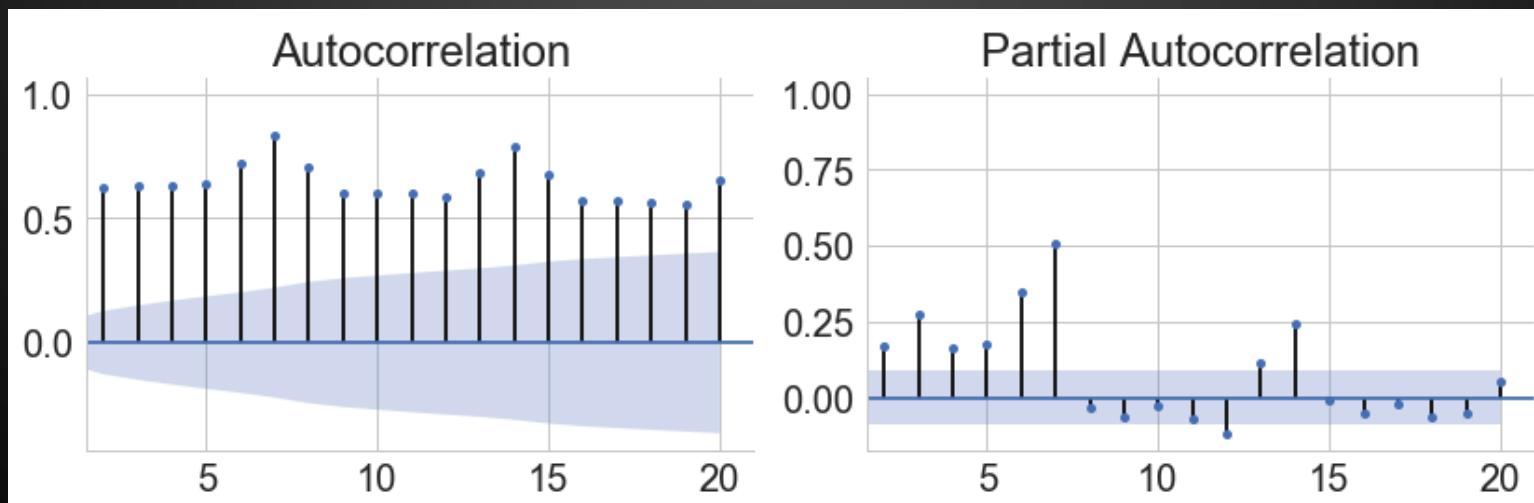
ACF being so high for so long  
suggests Integrated [ $d > 0$ ]

suggests Seasonal [ $s=7$ ]

as does a test for Stationarity,  
Adjusted Dickey-Fuller Test

Try:

$$(p, d, q) (P, D, Q)_s = (0, 1, 0) (1, 0, 1)_7$$



# Forecasting with ARIMA

$$(p, d, q) (P, D, Q)_s = (0, 1, 0) (1, 0, 1)_7$$



```
import statsmodels.tsa.api as smt

mod = smt.SARIMAX(nhpp_bc, order=(0, 1, 0), seasonal_order=(1, 0, 1, 7))
res = mod.fit()
```

# Forecasting with ARIMA

$$(p,d,q)(P,D,Q)_s = (0,1,0)(1,0,1)_7$$

6	res.summary()					
Statespace Model Results						
<b>Dep. Variable:</b>	transactions					
<b>No. Observations:</b>	500					
<b>Model:</b>	SARIMAX(0, 1, 0)x(1, 0, 1, 7)					
<b>Log Likelihood</b>	-694.806					
<b>Date:</b>	Wed, 06 Dec 2017					
<b>AIC</b>	1395.612					
<b>Time:</b>	23:43:00					
<b>BIC</b>	1408.256					
<b>Sample:</b>	03-02-2017 - 07-14-2018					
<b>HQIC</b>	1400.574					
<b>Covariance Type:</b>	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.S.L7	0.9994	0.001	940.764	0.000	0.997	1.001
ma.S.L7	-0.9669	0.026	-37.099	0.000	-1.018	-0.916
sigma2	0.9144	0.059	15.538	0.000	0.799	1.030
<b>Ljung-Box (Q):</b>	189.81	<b>Jarque-Bera (JB):</b>	0.01			
<b>Prob(Q):</b>	0.00	<b>Prob(JB):</b>	1.00			
<b>Heteroskedasticity (H):</b>	0.75	<b>Skew:</b>	0.00			
<b>Prob(H) (two-sided):</b>	0.06	<b>Kurtosis:</b>	3.02			

# Forecasting with ARIMA

$$(p, d, q) (P, D, Q)_s = (0, 1, 0) (1, 0, 1)_7$$

6   res.summary()						
Statespace Model Results						
<b>Dep. Variable:</b>		transactions		<b>No. Observations:</b>	500	
<b>Model:</b>		SARIMAX(0, 1, 0)x(1, 0, 1, 7)		<b>Log Likelihood</b>	-694.806	
<b>Date:</b>		Wed, 06 Dec 2017		<b>AIC</b>	1395.612	
<b>Time:</b>		23:43:00		<b>BIC</b>	1408.256	
<b>Sample:</b>		03-02-2017 - 07-14-2018		<b>HQIC</b>	1400.574	
<b>Covariance Type:</b> opg						
	<b>coef</b>	<b>std err</b>	<b>z</b>	<b>P&gt; z </b>	<b>[0.025</b>	<b>0.975]</b>
ar.S.L7	0.9994	0.001	940.764	0.000	0.997	1.001
ma.S.L7	-0.9669	0.026	-37.099	0.000	-1.018	-0.916
sigma2	0.9144	0.059	15.538	0.000	0.799	1.030
<b>Ljung-Box (Q):</b> 189.81 <b>Jarque-Bera (JB):</b> 0.01						
<b>Prob(Q):</b> 0.00			<b>Prob(JB):</b> 1.00			
<b>Heteroskedasticity (H):</b> 0.75			<b>Skew:</b> 0.00			
<b>Prob(H) (two-sided):</b> 0.06			<b>Kurtosis:</b> 3.02			

AIC

measures goodness-of-fit but punishes for using more parameters.  
Smaller AIC implies a better model.

# Forecasting with ARIMA

$$(p, d, q) (P, D, Q)_s = (0, 1, 0) (1, 0, 1)_7$$

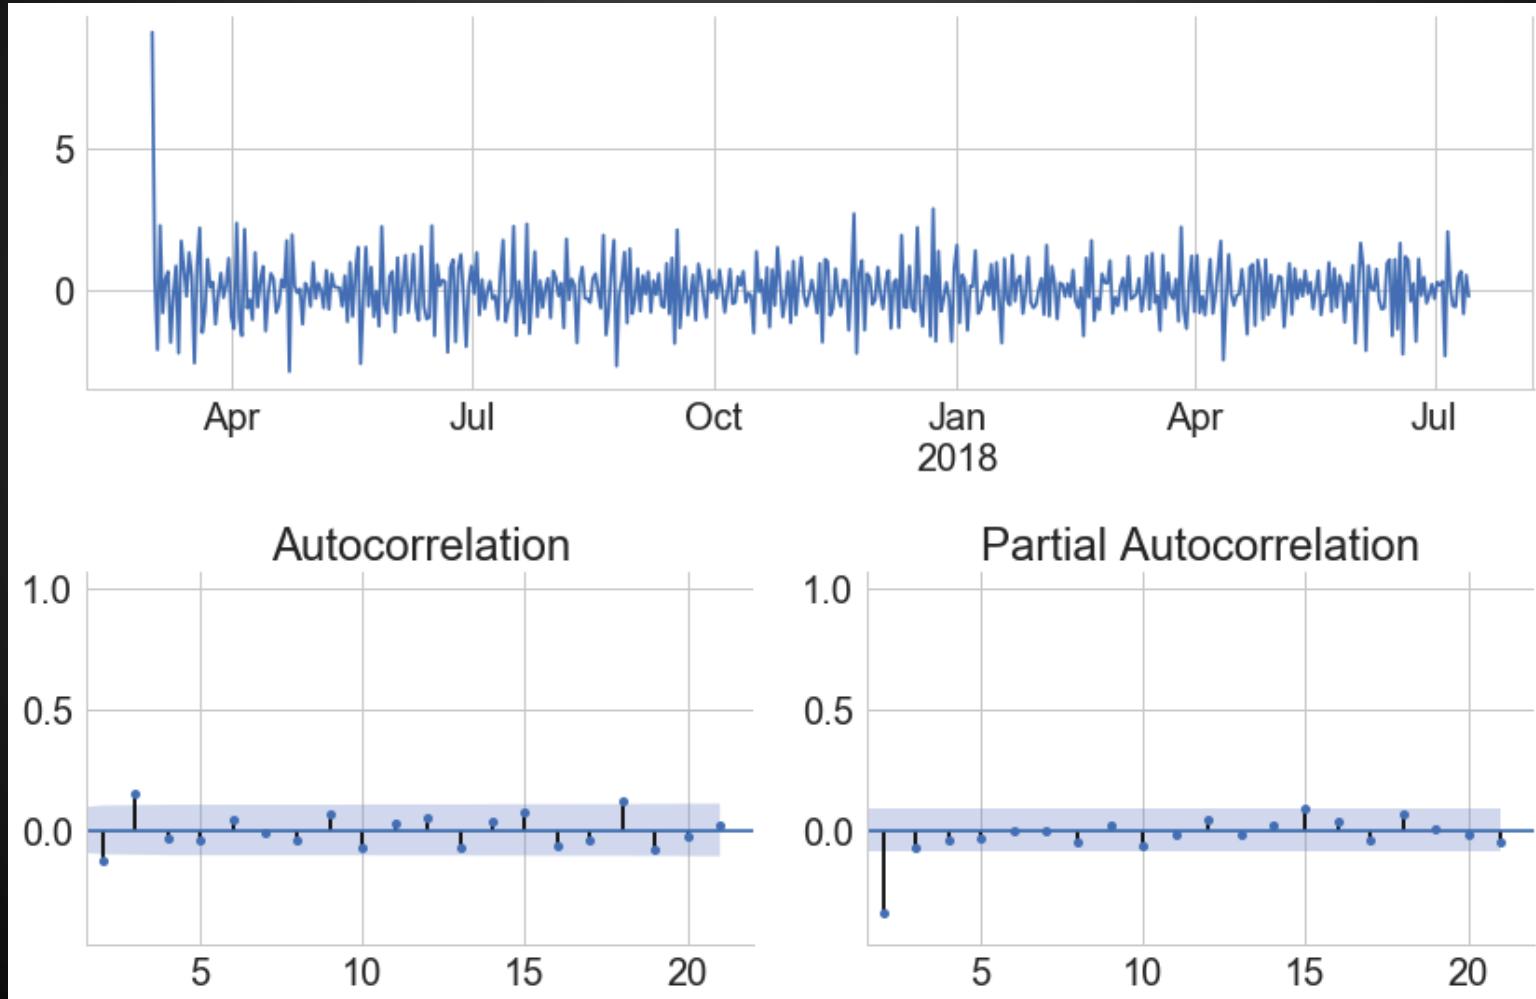
6   res.summary()																																		
Statespace Model Results																																		
<b>Dep. Variable:</b>		transactions		<b>No. Observations:</b>	500																													
<b>Model:</b>		SARIMAX(0, 1, 0)x(1, 0, 1, 7)		<b>Log Likelihood</b>	-694.806																													
<b>Date:</b>		Wed, 06 Dec 2017		<b>AIC</b>	1395.612																													
<b>Time:</b>		23:43:00		<b>BIC</b>	1408.256																													
<b>Sample:</b>		03-02-2017 - 07-14-2018		<b>HQIC</b>	1400.574																													
<b>Covariance Type:</b> opg																																		
<table><thead><tr><th></th><th>coef</th><th>std err</th><th>z</th><th>P&gt; z </th><th>[0.025</th><th>0.975]</th></tr></thead><tbody><tr><td>ar.S.L7</td><td>0.9994</td><td>0.001</td><td>940.764</td><td>0.000</td><td>0.997</td><td>1.001</td></tr><tr><td>ma.S.L7</td><td>-0.9669</td><td>0.026</td><td>-37.099</td><td>0.000</td><td>-1.018</td><td>-0.916</td></tr><tr><td>sigma2</td><td>0.9144</td><td>0.059</td><td>15.538</td><td>0.000</td><td>0.799</td><td>1.030</td></tr></tbody></table>								coef	std err	z	P> z	[0.025	0.975]	ar.S.L7	0.9994	0.001	940.764	0.000	0.997	1.001	ma.S.L7	-0.9669	0.026	-37.099	0.000	-1.018	-0.916	sigma2	0.9144	0.059	15.538	0.000	0.799	1.030
	coef	std err	z	P> z	[0.025	0.975]																												
ar.S.L7	0.9994	0.001	940.764	0.000	0.997	1.001																												
ma.S.L7	-0.9669	0.026	-37.099	0.000	-1.018	-0.916																												
sigma2	0.9144	0.059	15.538	0.000	0.799	1.030																												
<b>Ljung-Box (Q):</b> 189.81 <b>Jarque-Bera (JB):</b> 0.01																																		
<b>Prob(Q):</b> 0.00			<b>Prob(JB):</b> 1.00																															
<b>Heteroskedasticity (H):</b> 0.75			<b>Skew:</b> 0.00																															
<b>Prob(H) (two-sided):</b> 0.06			<b>Kurtosis:</b> 3.02																															

$\Phi_1$  →  
 $\Theta_1$  →

**AIC**  
measures goodness-of-fit but punishes for using more parameters.  
Smaller AIC implies a better model.

# Forecasting with ARIMA

$$(p, d, q) (P, D, Q)_s = (0, 1, 0) (1, 0, 1)_7$$

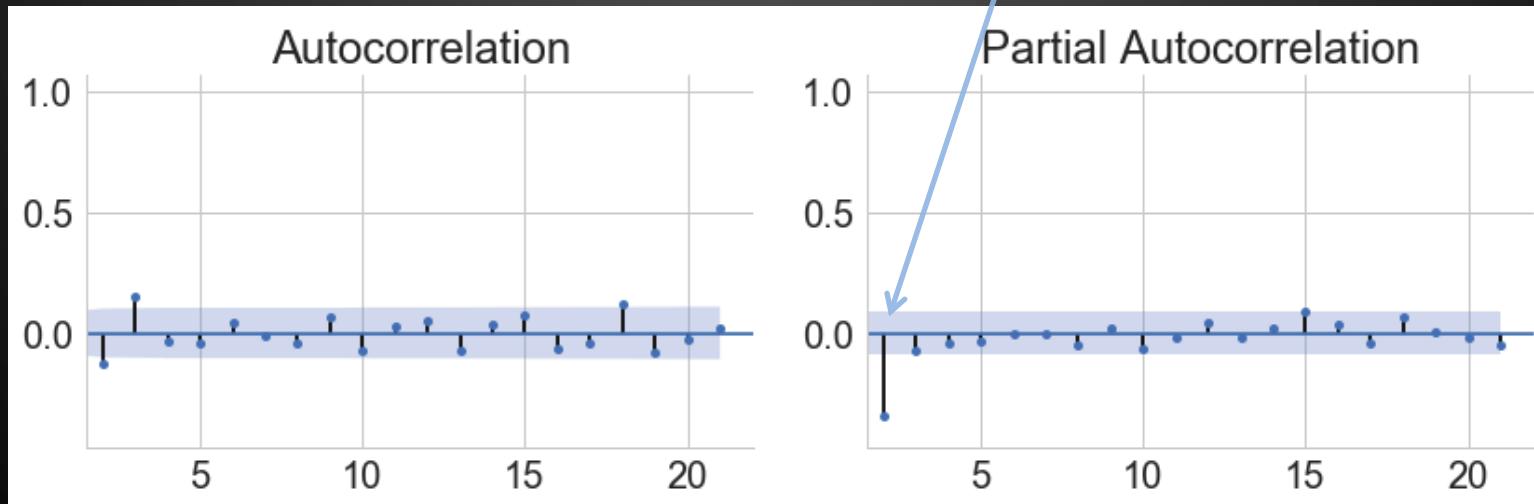


# Forecasting with ARIMA

$$(p, d, q) (P, D, Q)_s = (0, 1, 0) (1, 0, 1)_7$$

PACF being non-zero at 1  
but falling quickly after  
suggests Moving Average [q=1]

MA(1) is memory-less,  
covariance for steps > 1  
should be 0



# Forecasting with ARIMA

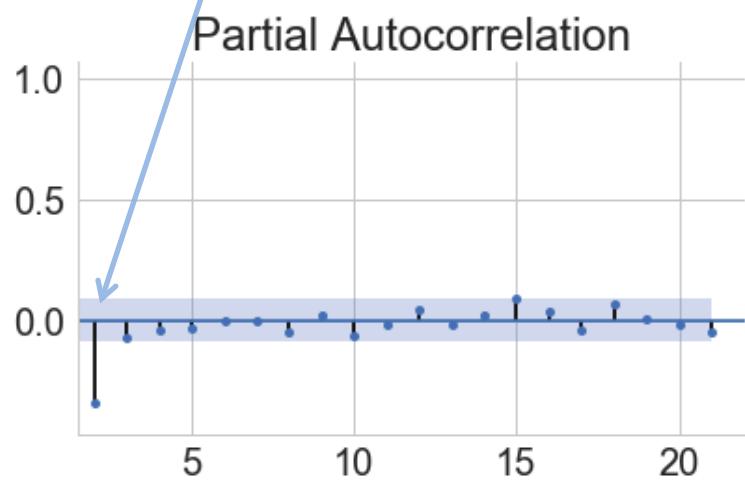
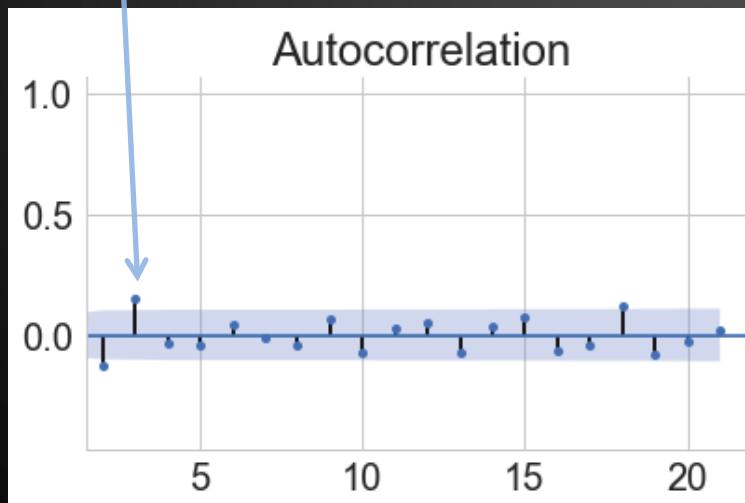
$$(p, d, q) (P, D, Q)_s = (0, 1, 0) (1, 0, 1)_7$$

ACF being non-zero above 1 suggests Autoregression [p>0]

covariance for steps > 0 decays as  $\phi^k$

PACF being non-zero at 1 but falling quickly after suggests Moving Average [q=1]

MA(1) is memory-less, covariance for steps > 1 should be 0



# Forecasting with ARIMA

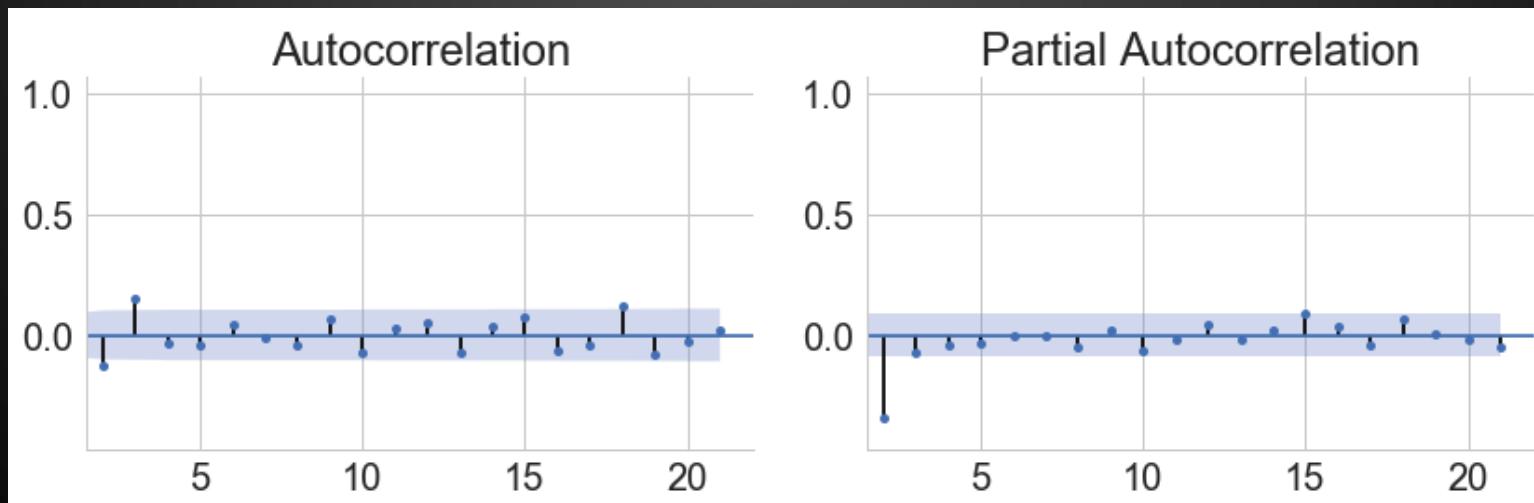
$$(p, d, q) (P, D, Q)_s = (0, 1, 0) (1, 0, 1)_7$$

ACF being non-zero above 1  
suggests Autoregression [p>0]

PACF being non-zero at 1  
but falling quickly after  
suggests Moving Average [q=1]

Try:

$$(p, d, q) (P, D, Q)_s = (2, 1, 1) (1, 0, 1)_7$$



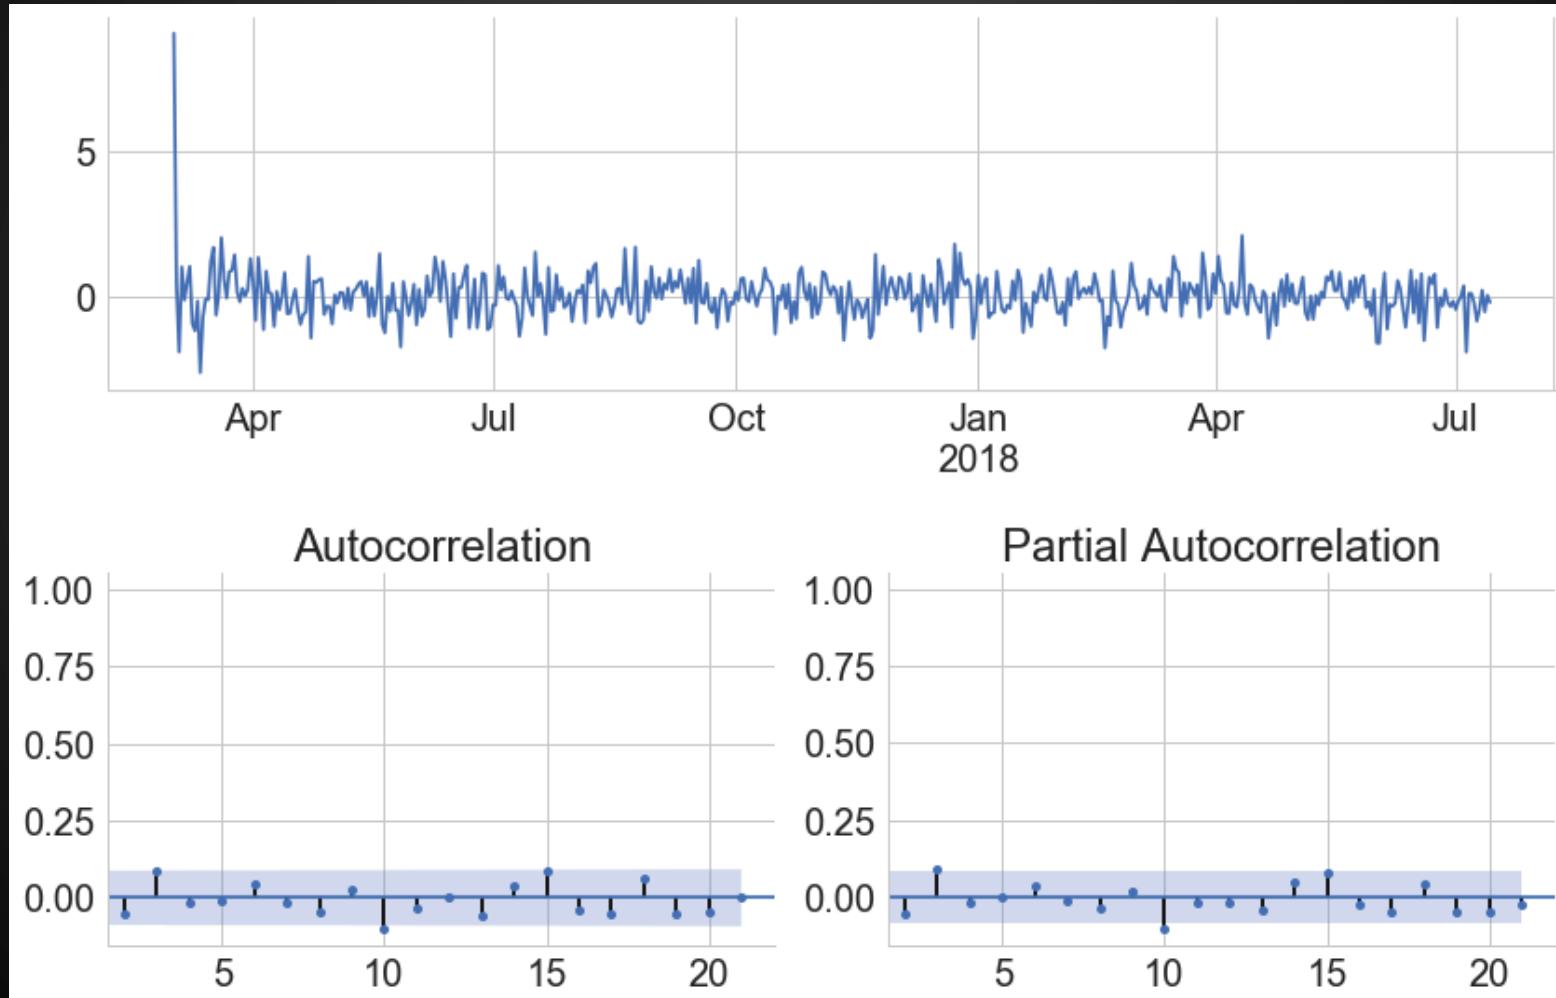
# Forecasting with ARIMA

$$(p, d, q) (P, D, Q)_s = (2, 1, 1) (1, 0, 1)_7$$

Statespace Model Results						
<b>Dep. Variable:</b>	transactions		<b>No. Observations:</b>	500		
<b>Model:</b>	SARIMAX(2, 1, 1)x(1, 0, 1, 7)			<b>Log Likelihood</b>	-514.523	
<b>Date:</b>	Wed, 06 Dec 2017			<b>AIC</b>	1041.046	
<b>Time:</b>	23:56:32			<b>BIC</b>	1066.333	
<b>Sample:</b>	03-02-2017 - 07-14-2018			<b>HQIC</b>	1050.968	
<b>Covariance Type:</b> opg						
	<b>coef</b>	<b>std err</b>	<b>z</b>	<b>P&gt; z </b>	<b>[0.025</b>	<b>0.975]</b>
ar.L1	-0.0629	0.046	-1.375	0.169	-0.152	0.027
ar.L2	-0.0986	0.049	-2.031	0.042	-0.194	-0.003
ma.L1	-0.9662	0.013	-75.824	0.000	-0.991	-0.941
ar.S.L7	0.9996	0.001	1602.365	0.000	0.998	1.001
ma.S.L7	-0.9670	0.025	-38.726	0.000	-1.016	-0.918
sigma2	0.4381	0.027	16.152	0.000	0.385	0.491
<b>Ljung-Box (Q):</b> 36.22 <b>Jarque-Bera (JB):</b> 0.71						
<b>Prob(Q):</b> 0.64			<b>Prob(JB):</b> 0.70			
<b>Heteroskedasticity (H):</b> 0.83			<b>Skew:</b> -0.04			
<b>Prob(H) (two-sided):</b> 0.24			<b>Kurtosis:</b> 3.16			

# Forecasting with ARIMA

$$(p, d, q) (P, D, Q)_s = (2, 1, 1) (1, 0, 1)_7$$



# Forecasting with ARIMA

gobear

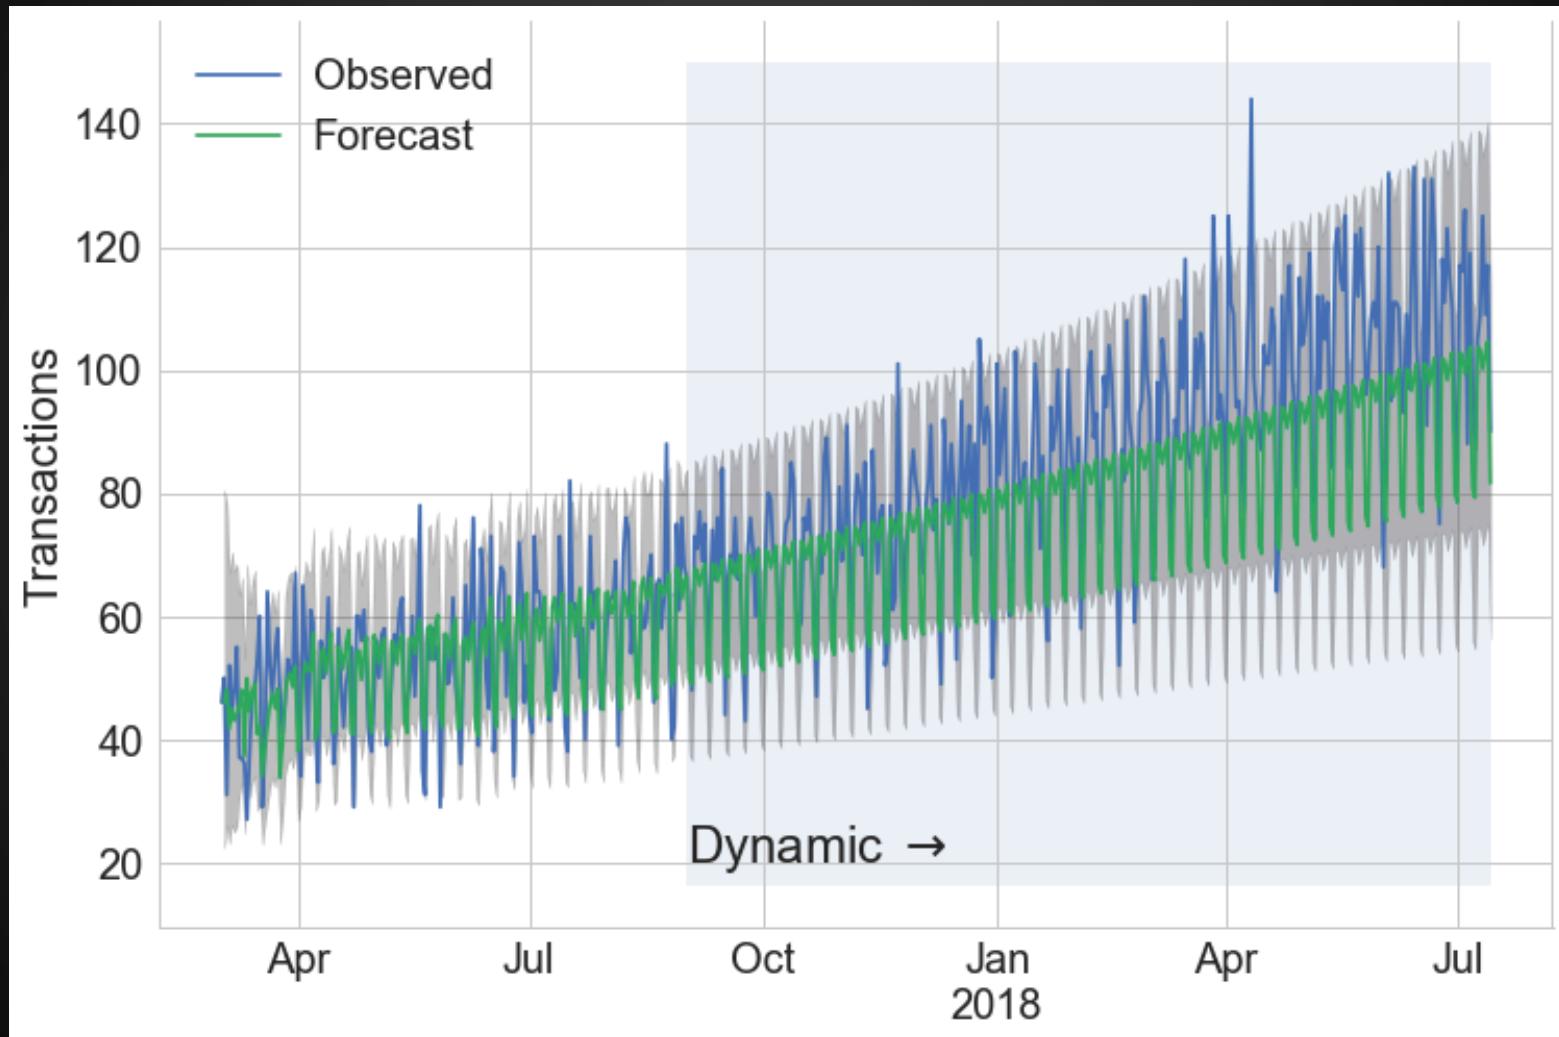
Y



# Forecasting with ARIMA

```
1 # Get the predictions, and forecast beyond dynamic_start
2 dynamic_start = '2017-09-01'
3 pred = res.get_prediction(start='2017-03-03', end=max(nhpp.index), dynamic=dynamic_start)
4
5 # Undo the Box-Cox transformation
6 predictions = (λ_bc * pred.predicted_mean + 1) ** (1/ λ_bc)
7 predictions_bounds = (λ_bc * pred.conf_int() + 1) ** (1 / λ_bc)
8
9 # Plot the time series and the forecast
10 ax = nhpp.plot(label='Observed')
11 predictions.plot(ax=ax, label='Forecast')
12 ax.fill_between(predictions.index,
13                 predictions_bounds.iloc[:, 0],
14                 predictions_bounds.iloc[:, 1], color='k', alpha=.25)
15
16 # Highlight the forecast area
17 ax.fill_betweenx(ax.get_ylim(), pd.Timestamp(dynamic_start), predictions.index[-1], alpha=0.1)
18 ax.annotate('Dynamic $\rightarrow$', (pd.Timestamp(dynamic_start), 20))
19 ax.set_ylabel('Transactions')
20 plt.legend()
21 ''
```

# Forecasting with ARIMA

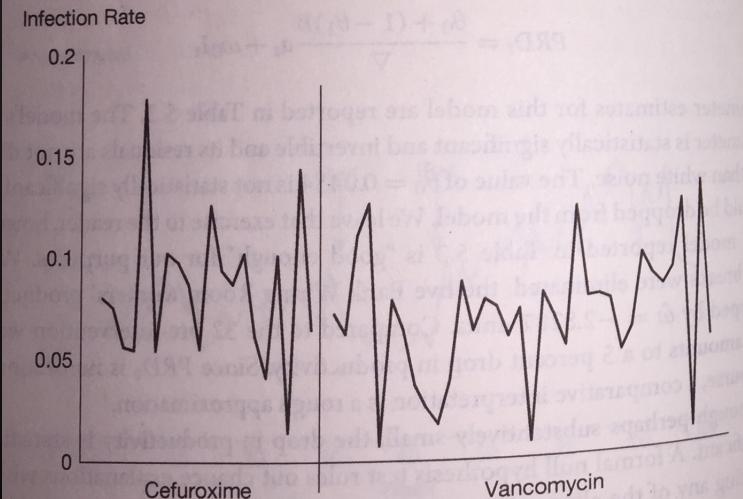


# Interrupted Time Series

Introduce variant at time  $t_{\text{start}}$  of effect size  $\beta$   
Basic idea : Fit ARIMA + effect size

for the next  $N_{\text{post}} = 35$  months, with propylactic vancomycin. Modeling time series will be especially difficult. The inherent difficulty of identifying  $t_{\text{start}}$  from a short time series is aggravated in this instance by the unruled appearance of this time series and the possibility of a large impact. We start the alternative ARIMA modeling strategy with the simplest possible model:

$$SSI_t = \theta_0 + a_t + \omega_0 I_t$$

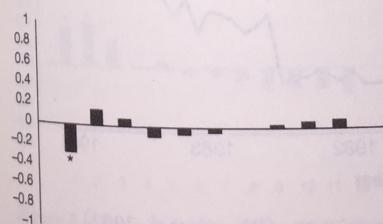


$$\begin{aligned} SSI_t &= 0.084 - 0.014I_t + \frac{a_t}{1 + .27B} \\ &= 0.084 - 0.014I_t + \sum_{k=0}^{\infty} 0.27^k a_t \\ &= 0.084 - 0.014I_t + a_t + 0.27a_{t-1} + 0.07a_{t-2} + 0.02a_{t-3} + \dots \end{aligned}$$

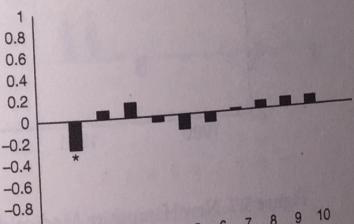
Table 5.4. SURGICAL SITE INFECTION RATES, ESTIMATION AND DIAGNOSIS

	Estimate	SE	t	Diagnosis	RSE
Model 1	$\hat{\theta}_0$	.0833	.0075	11.04	0.03373
	$\hat{\omega}_0$	-.0141	.0095	-1.49	$Q_{10} = 7.4$ $Z_{KS} = 0.805$
Model 2	$\hat{\theta}_0$	.0836	.0059	14.11	0.03277
	$\hat{\phi}_1$	-.2710	.1312	-2.07	$Q_{10} = 3.9$ $Z_{KS} = 0.805$
	$\hat{\omega}_0$	-.0144	.0074	-1.95	

Sample ACF for  $\hat{a}_t$



Sample PACF for  $\hat{a}_t$



source: McCleary, et al. "Design and Analysis of Time Series Experiments"

# Interrupted Time Series

Start with a simple time series, for illustration

```
df = poisson_process.nhpp(  
    λ=100, n=10000, t0='2017-03-01', freq='1D',  
    thinning_components=[{'function': poisson_process._trend,  
        'params': {'percent_increase': 1.0, 'period': 365}}]).to_frame()
```

# Interrupted Time Series

Add a small effect size at some start time

```
df = poisson_process.nhpp(  
    λ=100, n=10000, t0='2017-03-01', freq='1D',  
    thinning_components=[{'function': poisson_process._trend,  
        'params': {'percent_increase': 1.0, 'period': 365}}]).to_frame()  
  
treatment_start = '2017-05-01'  
effect_size = 0.1  
df.loc[treatment_start:, 'response'] = (  
    np.random.normal(loc=effect_size, scale=effect_size / 2, size=len(df[treatment_start:])))  
df['transactions_with_response'] = (df['transactions'] * (1 + df['response'].fillna(0))).astype(int)
```

# Interrupted Time Series

Find the ARIMA model that fits best

```
df = poisson_process.nhpp(  
    λ=100, n=10000, t0='2017-03-01', freq='1D',  
    thinning_components=[{'function': poisson_process._trend,  
        'params': {'percent_increase': 1.0, 'period': 365}}]).to_frame()  
  
treatment_start = '2017-05-01'  
effect_size = 0.1  
df.loc[treatment_start:, 'response'] = (  
    np.random.normal(loc=effect_size, scale=effect_size / 2, size=len(df[treatment_start:])))  
df['transactions_with_response'] = (df['transactions'] * (1 + df['response'].fillna(0))).astype(int)  
  
mod = smt.SARIMAX(df.loc[:treatment_start, 'transactions'],  
                    order=(0, 1, 1),  
                    seasonal_order=(0, 0, 0, 0))  
res = mod.fit()
```

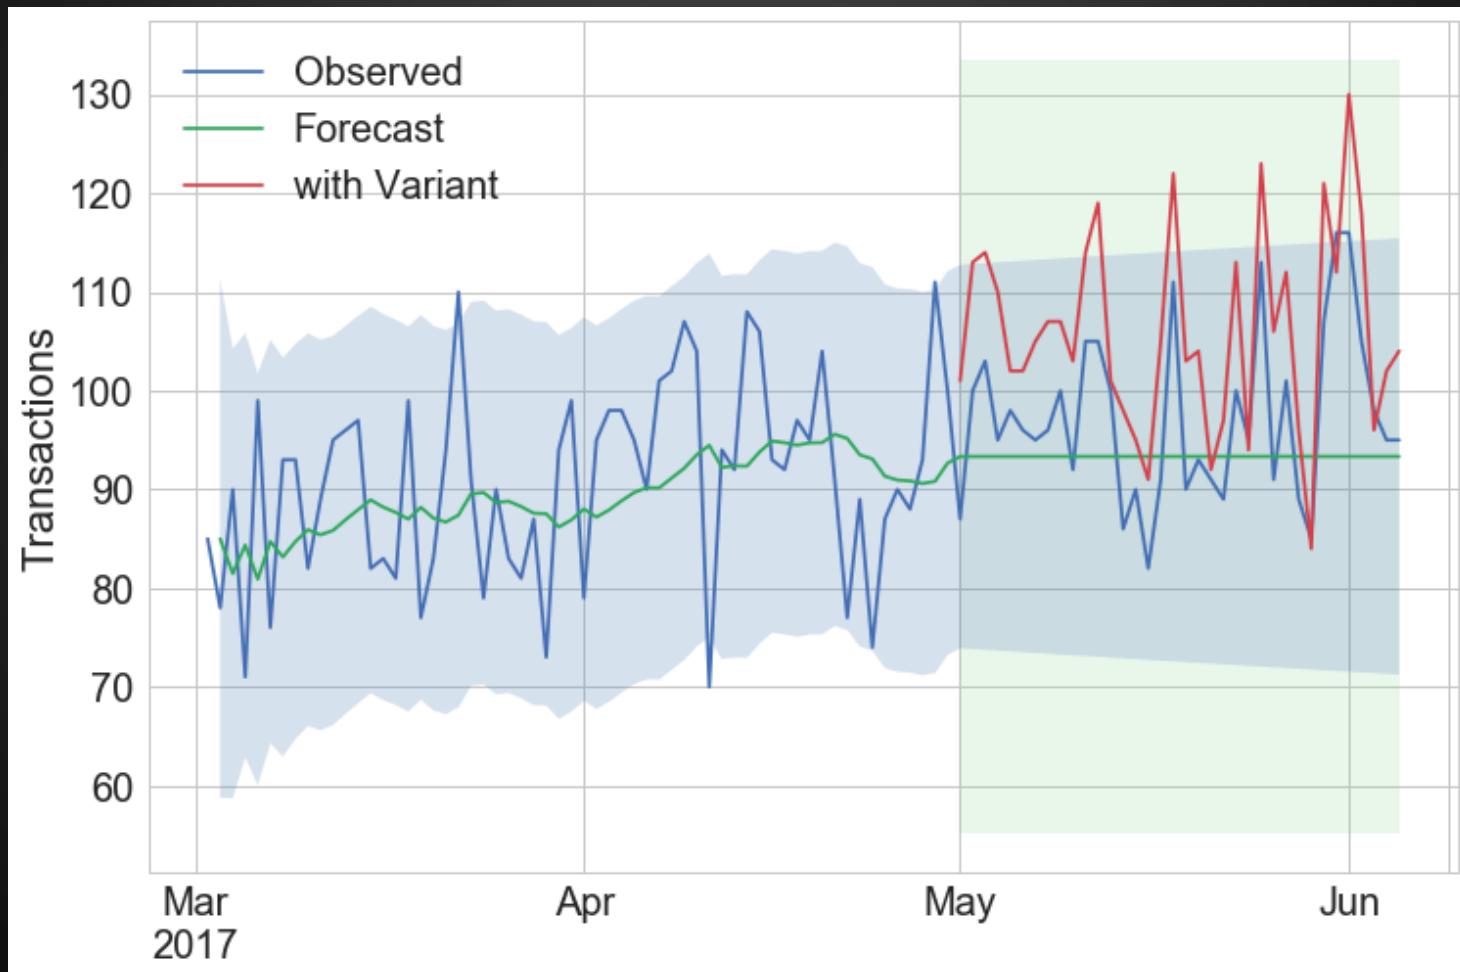
# Interrupted Time Series

Get the predictions and confidence bands

```
df = poisson_process.nhpp(  
    λ=100, n=10000, t0='2017-03-01', freq='1D',  
    thinning_components=[{'function': poisson_process._trend,  
        'params': {'percent_increase': 1.0, 'period': 365}}]).to_frame()  
  
treatment_start = '2017-05-01'  
effect_size = 0.1  
df.loc[treatment_start:, 'response'] = (  
    np.random.normal(loc=effect_size, scale=effect_size / 2, size=len(df[treatment_start:])))  
df['transactions_with_response'] = (df['transactions'] * (1 + df['response'].fillna(0))).astype(int)  
  
mod = smt.SARIMAX(df.loc[:treatment_start, 'transactions'],  
                    order=(0, 1, 1),  
                    seasonal_order=(0, 0, 0, 0))  
res = mod.fit()  
  
pred = res.get_prediction(start='2017-03-03', end=max(df.index), dynamic=treatment_start)  
df['prediction'] = pred.predicted_mean  
df[['prediction_lower', 'prediction_upper']] = pred.conf_int()
```

# Interrupted Time Series

Our Data Set



# Interrupted Time Series

```
import pymc3 as pm
from theano import scan, shared
import theano.tensor as tt

def build_model(X, treatment_start, treatment_observations):
    time_seen = pd.to_datetime(treatment_start) + pd.DateOffset(treatment_observations - 1)
    y = shared(X[:time_seen].values)
    y_switch = shared(X[:time_seen].index < treatment_start)
    with pm.Model() as ilmal:
        σ = pm.HalfCauchy('σ', beta=2.)
        θ = pm.Normal('θ', 0., sd=2.)
        β = pm.Normal('β', 0., sd=2.)

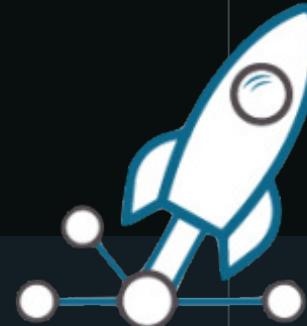
        y_adj = tt.switch(y_switch, y, y - tt.dot(y, β))

        # ARIMA (0, 1, 1)
        # -----
        # (1 - B) y[t] = (1 - θB) ε[t]
        # y[t] - y[t-1] = ε[t] - θ * ε[t-1]
        # ε[t] = y[t] - y[t-1] - θ * ε[t-1]
        def calc_next(y_lag1, y_lag0, ε, θ):
            return y_lag0 - y_lag1 - θ * ε

        # Initial noise guess -- let's just seed with 0
        ε0 = tt.zeros_like(y_adj)

        ε, _ = scan(fn=calc_next,
                    sequences=dict(input=y_adj, taps=[-1, 0]),
                    outputs_info=[ε0],
                    non_sequences=[θ])

        pm.Potential('like', pm.Normal.dist(0, sd=σ).logp(ε))
    return ilmal
```



# Interrupted Time Series

```
import pymc3 as pm
from theano import scan, shared
import theano.tensor as tt

def build_model(X, treatment_start, treatment_observations):
    time_seen = pd.to_datetime(treatment_start) + pd.DateOffset(treatment_observations - 1)
    y = shared(X[:time_seen].values)
    y_switch = shared(X[:time_seen].index < treatment_start)
    with pm.Model() as ilmal:
        σ = pm.HalfCauchy('σ', beta=2.)
        θ = pm.Normal('θ', 0., sd=2.)
        β = pm.Normal('β', 0., sd=2.)
        y_adj = tt.switch(y_switch, y, y - tt.dot(y, β))

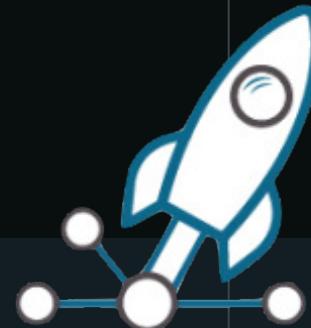
        # ARIMA (0, 1, 1)
        # -----
        # (1 - B) y[t] = (1 - θB) ε[t]
        # y[t] - y[t-1] = ε[t] - θ * ε[t-1]
        # ε[t] = y[t] - y[t-1] - θ * ε[t-1]
        def calc_next(y_lag1, y_lag0, ε, θ):
            return y_lag0 - y_lag1 - θ * ε

        # Initial noise guess -- let's just seed with 0
        ε0 = tt.zeros_like(y_adj)

        ε, _ = scan(fn=calc_next,
                    sequences=dict(input=y_adj, taps=[-1, 0]),
                    outputs_info=[ε0],
                    non_sequences=[θ])

        pm.Potential('like', pm.Normal.dist(0, sd=σ).logp(ε))
    return ilmal
```

Case statement:  
left = control  
right = variant  
[effect.size =  $\beta$  ]



# Interrupted Time Series

```
import pymc3 as pm
from theano import scan, shared
import theano.tensor as tt

def build_model(X, treatment_start, treatment_observations):
    time_seen = pd.to_datetime(treatment_start) + pd.DateOffset(treatment_observations - 1)
    y = shared(X[:time_seen].values)
    y_switch = shared(X[:time_seen].index < treatment_start)
    with pm.Model() as ilmal:
        σ = pm.HalfCauchy('σ', beta=2.)
        θ = pm.Normal('θ', 0., sd=2.)
        β = pm.Normal('β', 0., sd=2.)
        y_adj = tt.switch(y_switch, y, y - tt.dot(y, β))

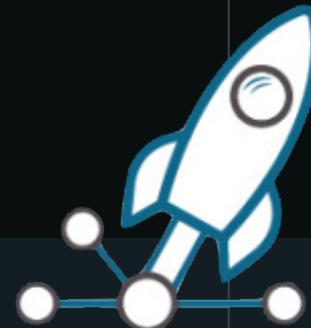
# ARIMA (0, 1, 1)
# -----
# (1 - B) y[t] = (1 - θB) ε[t]
# y[t] - y[t-1] = ε[t] - θ * ε[t-1]
# ε[t] = y[t] - y[t-1] - θ * ε[t-1]
def calc_next(y_lag1, y_lag0, ε, θ):
    return y_lag0 - y_lag1 - θ * ε

# Initial noise guess -- let's just seed with 0
ε0 = tt.zeros_like(y_adj)

ε, _ = scan(fn=calc_next,
            sequences=dict(input=y_adj, taps=[-1, 0]),
            outputs_info=[ε0],
            non_sequences=[θ])

pm.Potential('like', pm.Normal.dist(0, sd=σ).logp(ε))
return ilmal
```

Case statement:  
left = control  
right = variant  
[effect.size =  $\beta$  ]



ARIMA (0, 1, 1)

# Interrupted Time Series

```
import pymc3 as pm
from theano import scan, shared
import theano.tensor as tt

def build_model(X, treatment_start, treatment_observations):
    time_seen = pd.to_datetime(treatment_start) + pd.DateOffset(treatment_observations - 1)
    y = shared(X[:time_seen].values)
    y_switch = shared(X[:time_seen].index < treatment_start)
    with pm.Model() as ilmal:
        σ = pm.HalfCauchy('σ', beta=2.)
        θ = pm.Normal('θ', 0., sd=2.)
        β = pm.Normal('β', 0., sd=2.)
        y_adj = tt.switch(y_switch, y, y - tt.dot(y, β))

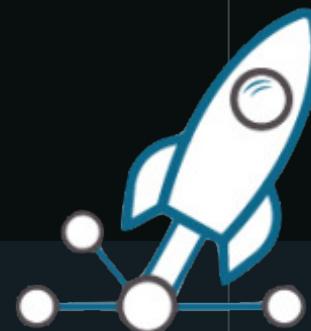
        # ARIMA (0, 1, 1)
        # -----
        # (1 - B) y[t] = (1 - θB) ε[t]
        # y[t] - y[t-1] = ε[t] - θ * ε[t-1]
        # ε[t] = y[t] - y[t-1] - θ * ε[t-1]
        def calc_next(y_lag1, y_lag0, ε, θ):
            return y_lag0 - y_lag1 - θ * ε

        # Initial noise guess -- let's just seed with 0
        ε0 = tt.zeros_like(y_adj)

        ε, _ = scan(fn=calc_next,
                    sequences=dict(input=y_adj, taps=[-1, 0]),
                    outputs_info=[ε0],
                    non_sequences=[θ])

        pm.Potential('like', pm.Normal.dist(0, sd=σ).logp(ε))
    return ilmal
```

Case statement:  
left = control  
right = variant  
[effect.size =  $\beta$  ]

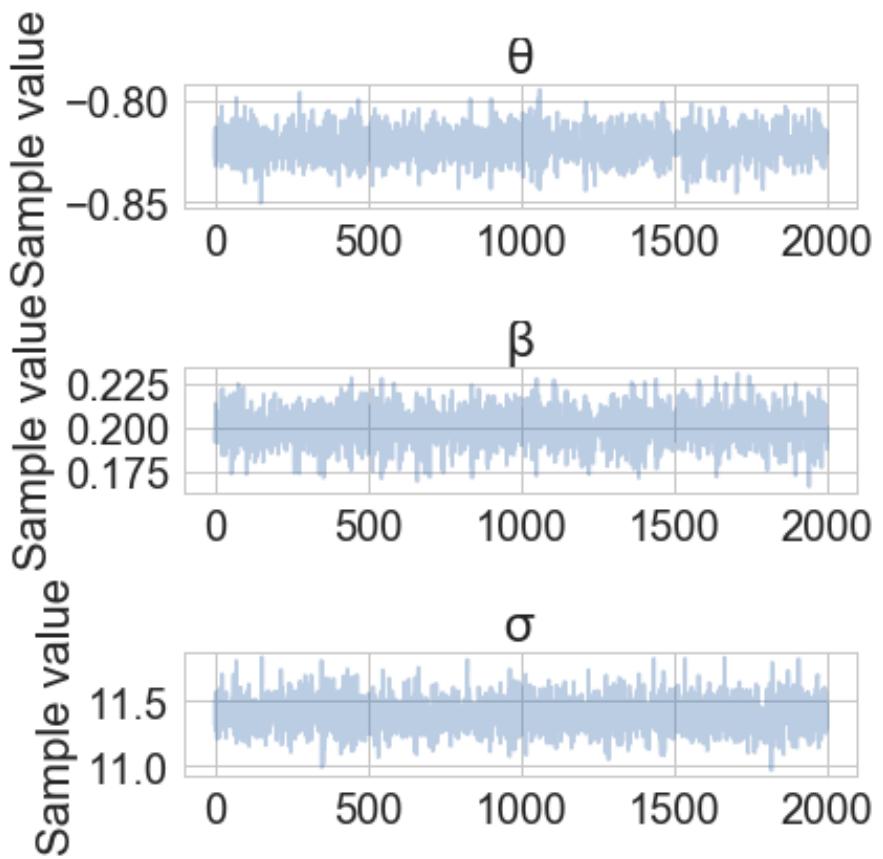
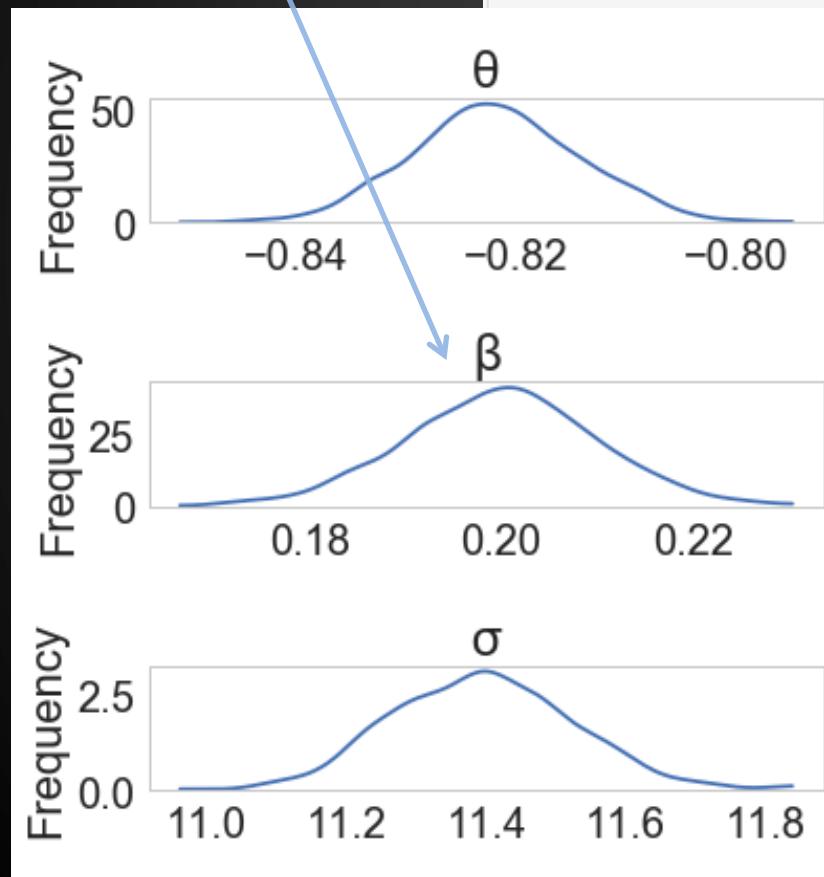


Looping in Theano

# Interrupted Time Series

effect size estimation

```
treatment_observations = 2
with build_model(df['transactions_with_response'],
                  treatment_start,
                  treatment_observations):
    trace = pm.sample(2000)
pm.traceplot(trace, combined=True)
```



# Interrupted Time Series



# Interrupted Time Series Experiments in Python

PhD, Physics

Data Scientist



SPOT  
HERO  
GRUBHUB



**Drew Fustin**  
Automaton Data  
*available for hire!*  
[drew@automatondata.com](mailto:drew@automatondata.com)

[github.com/drewfustin/talks\\_time-series-experiments](https://github.com/drewfustin/talks_time-series-experiments)

@drewfustin | PyData Ann Arbor | 2017-12-07